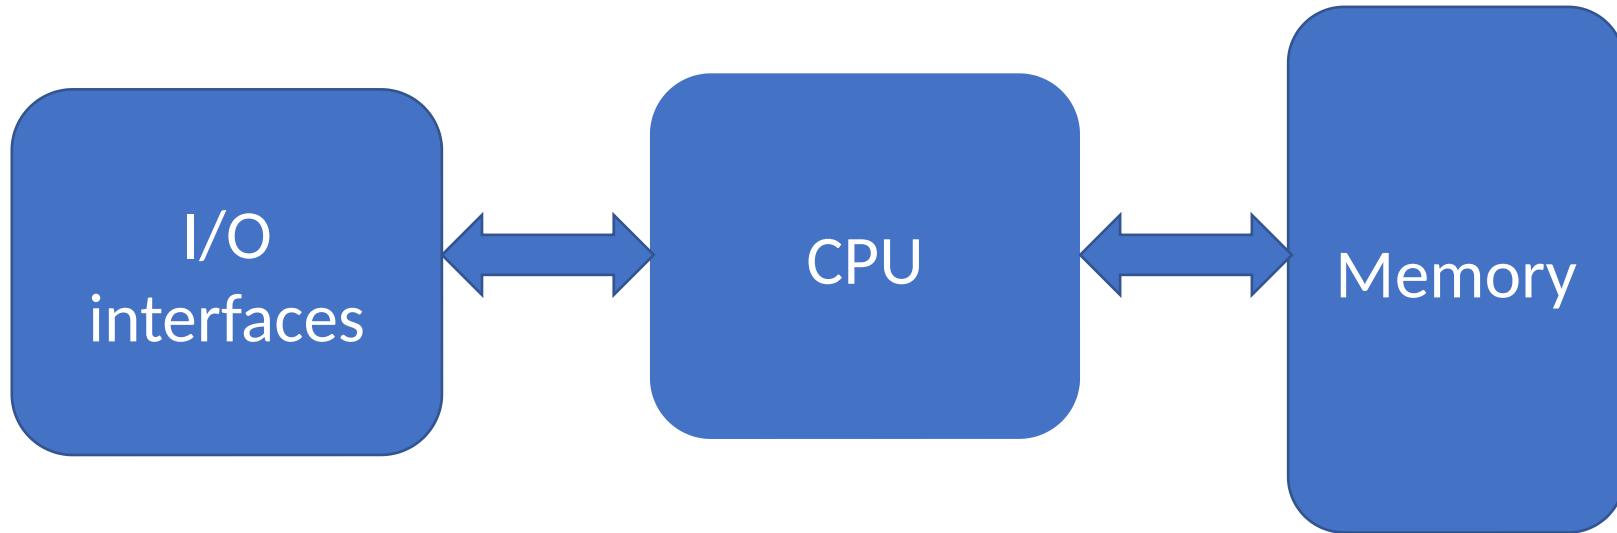


Multicore Systems

Eike Ritter
School of Computer Science
University of Birmingham

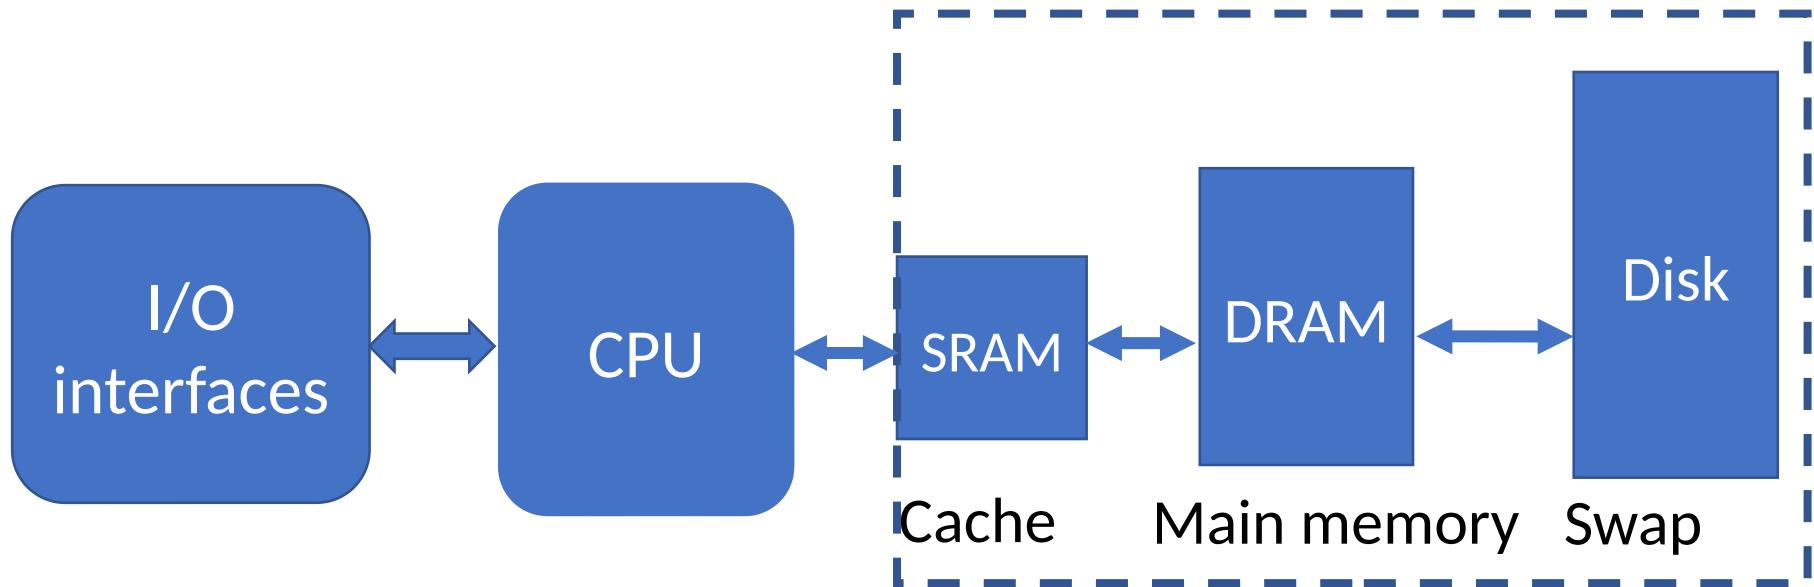
Recap: Simple von Neumann Computer



We started C programming assuming a simple von Neumann computation model.

- We executed program on a single CPU (or processor) step-by-step.
- We stored data and program instructions in the Memory.

Recap: Memory Hierarchy



We studied memory hierarchy from a programmer's perspective
→ Programs are still executed on a single CPU (or processor)
step-by-step but more efficiently

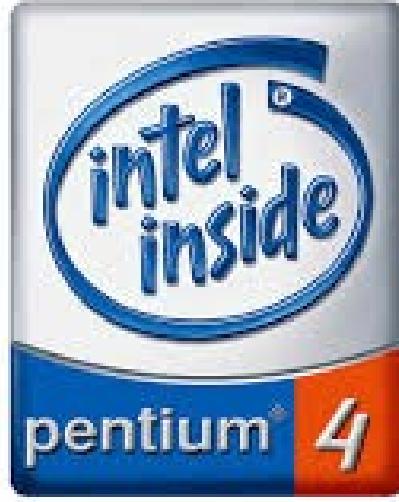


Produced	From May 7, 1997 to December 26, 2003 ^[1]
Common manufacturer(s)	Intel
Max. CPU clock rate	233 MHz to 450 MHz
FSB speeds	66 MHz to 100 MHz
Min. feature size	0.35 µm to 0.18 µm
Instruction set	IA-32, MMX



Produced	From February 26, 1999 to May 18, 2007 ^[1]
Common manufacturer(s)	Intel
Max. CPU clock rate	450 MHz to 1.4 GHz
FSB speeds	100 MHz to 133 MHz
Min. feature size	0.25 µm to 0.13 µm
Instruction set	IA-32, MMX, SSE

Clock frequency almost doubled w.r.t Pentium II



Produced	From November 20, 2000 to August 8, 2008
Max. CPU clock rate	1.3 GHz to 3.8 GHz
FSB speeds	400 MT/s to 1066 MT/s
Instruction set	x86 (i386), x86-64 (only some chips), MMX, SSE, SSE2, SSE3

Clock frequency almost doubled w.r.t Pentium III

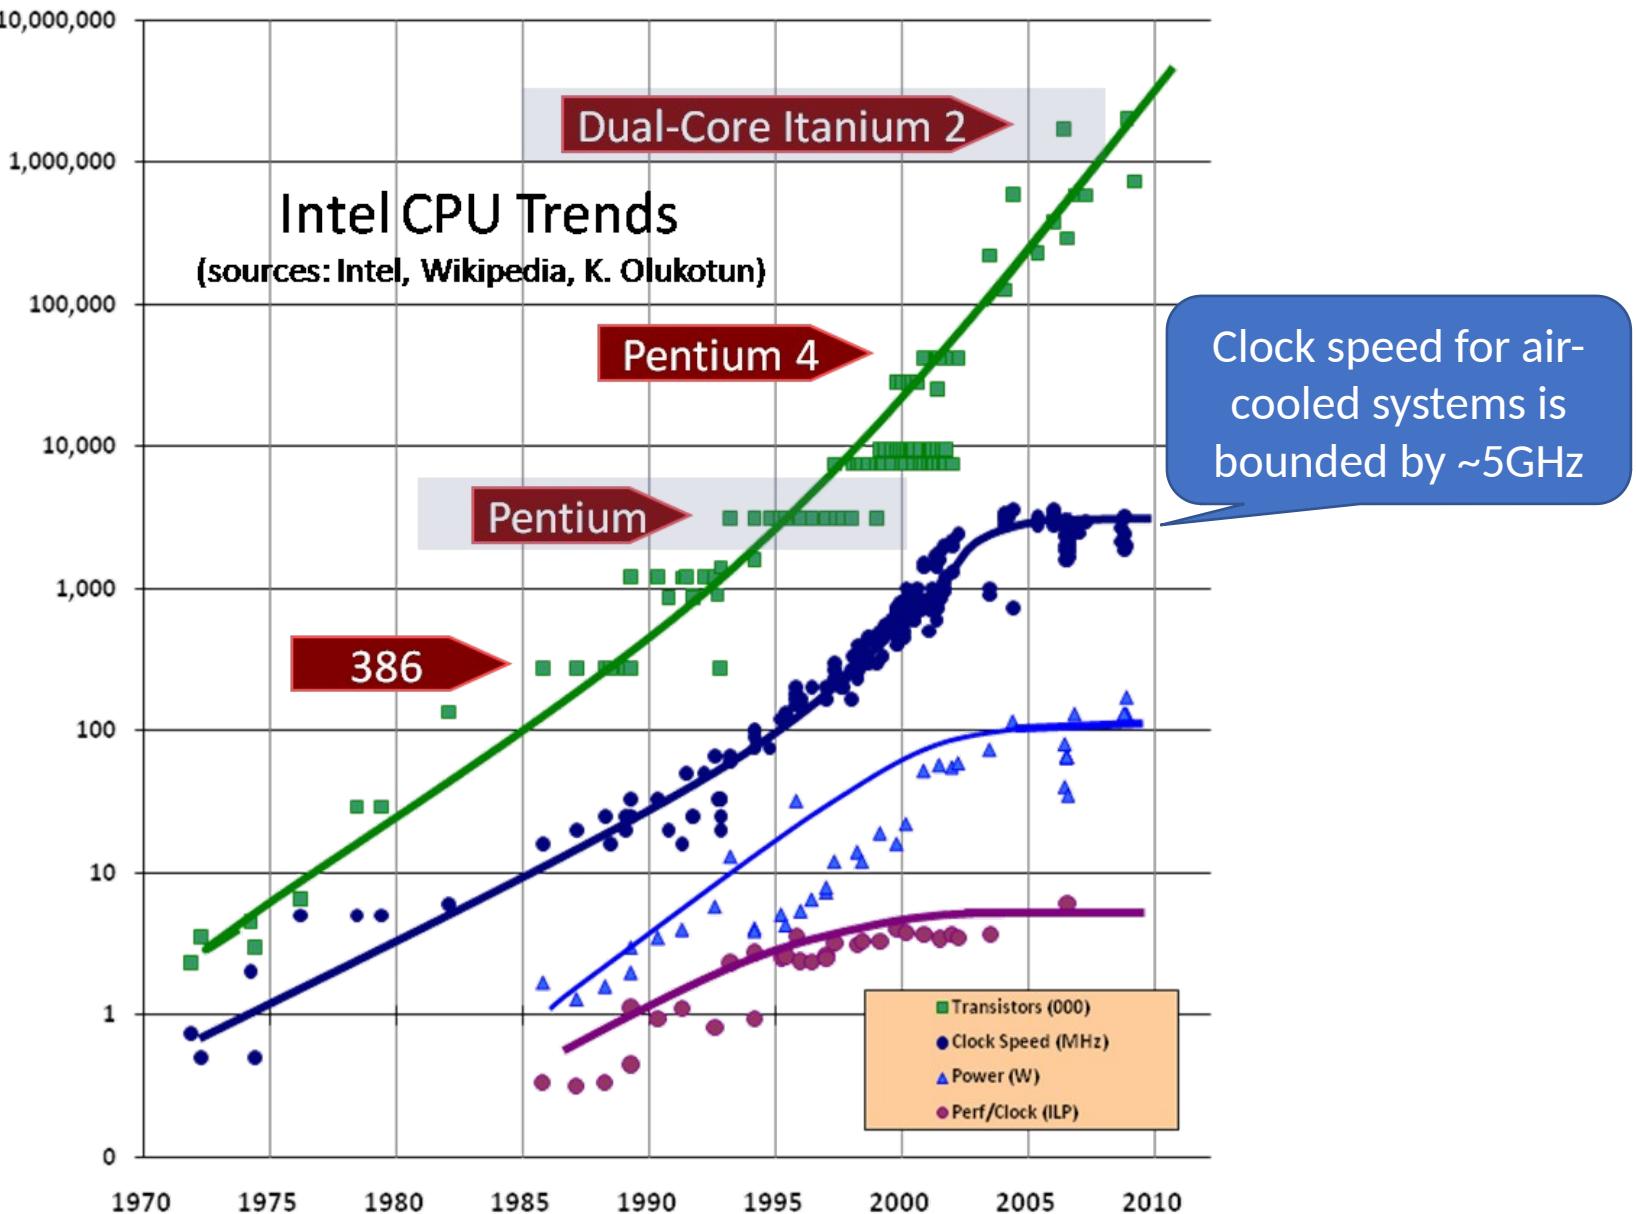
What is the clock frequency in 2020?



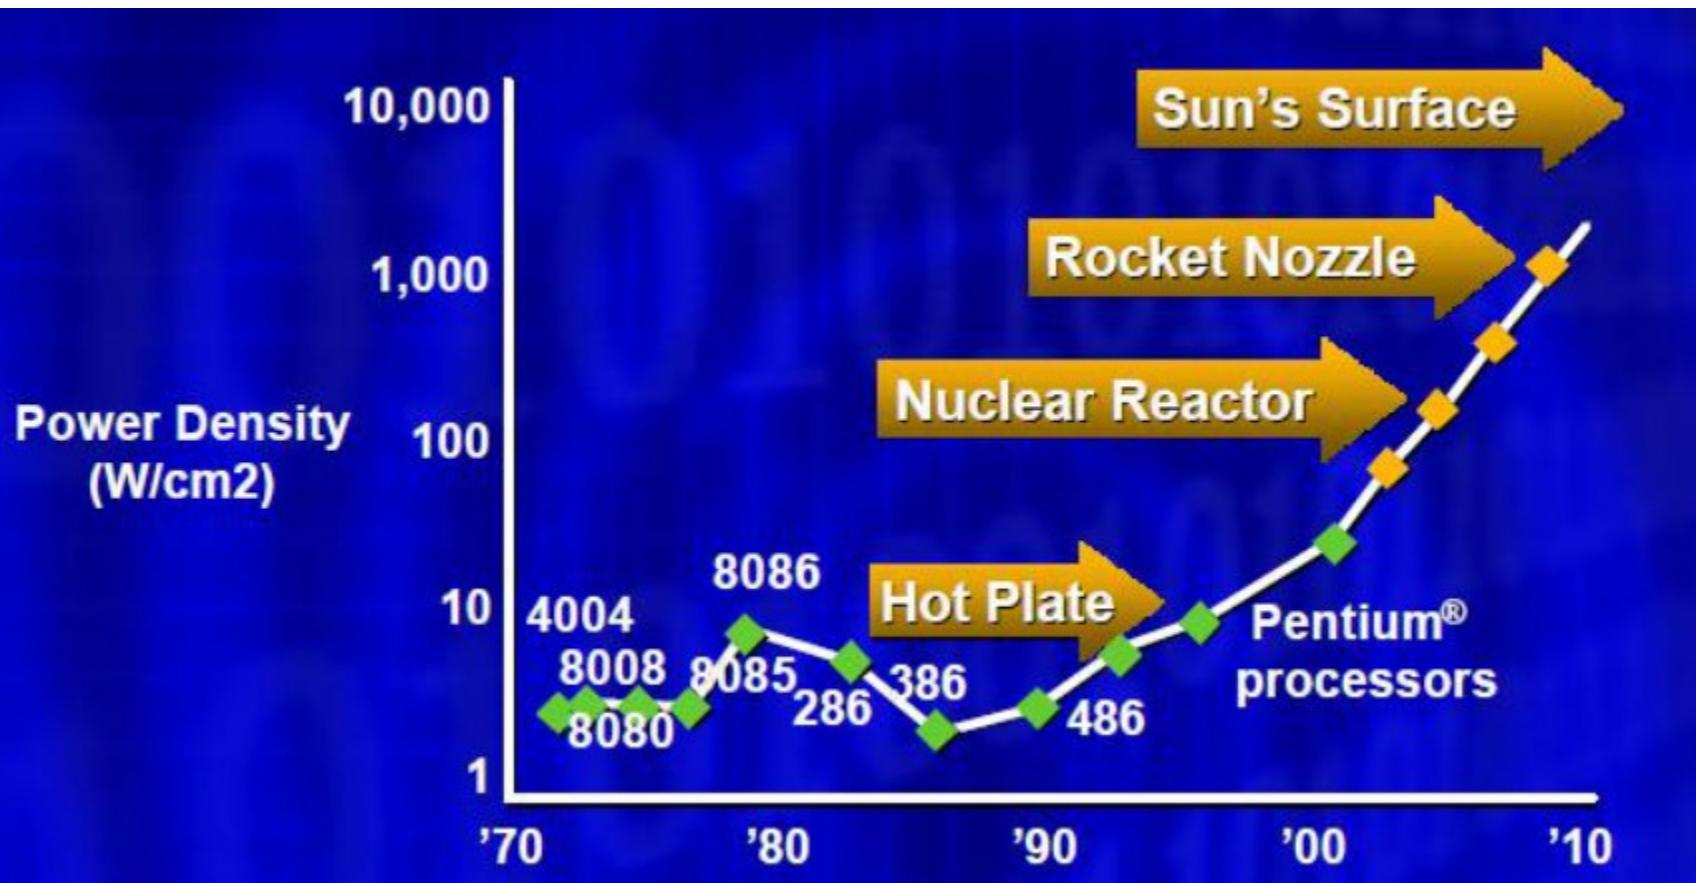
Intel® Core™ i7-10510U Processor (8M Cache, Up to 4.80 GHz)

- 8 MB Intel® Smart Cache Cache
- 4 Cores
- 8 Threads
- 4.90 GHz Max Turbo Frequency
- U - Ultra-low power
- 10th Generation

Little increase in clock frequency after 10 years since Pentium 4



Source: Herb Sutter “The Free Lunch Is Over: A Fundamental Turn Toward Concurrency in Software” available at <http://www.gotw.ca/publications/concurrency-ddj.htm>



Source: Patrick Gelsinger, Intel Developer's Forum, Intel Corporation, 2004.

Applications are demanding more resources!

Alternative solution:

- Put many processing cores on the microprocessor chip.
- The number of cores doubles with each generation.

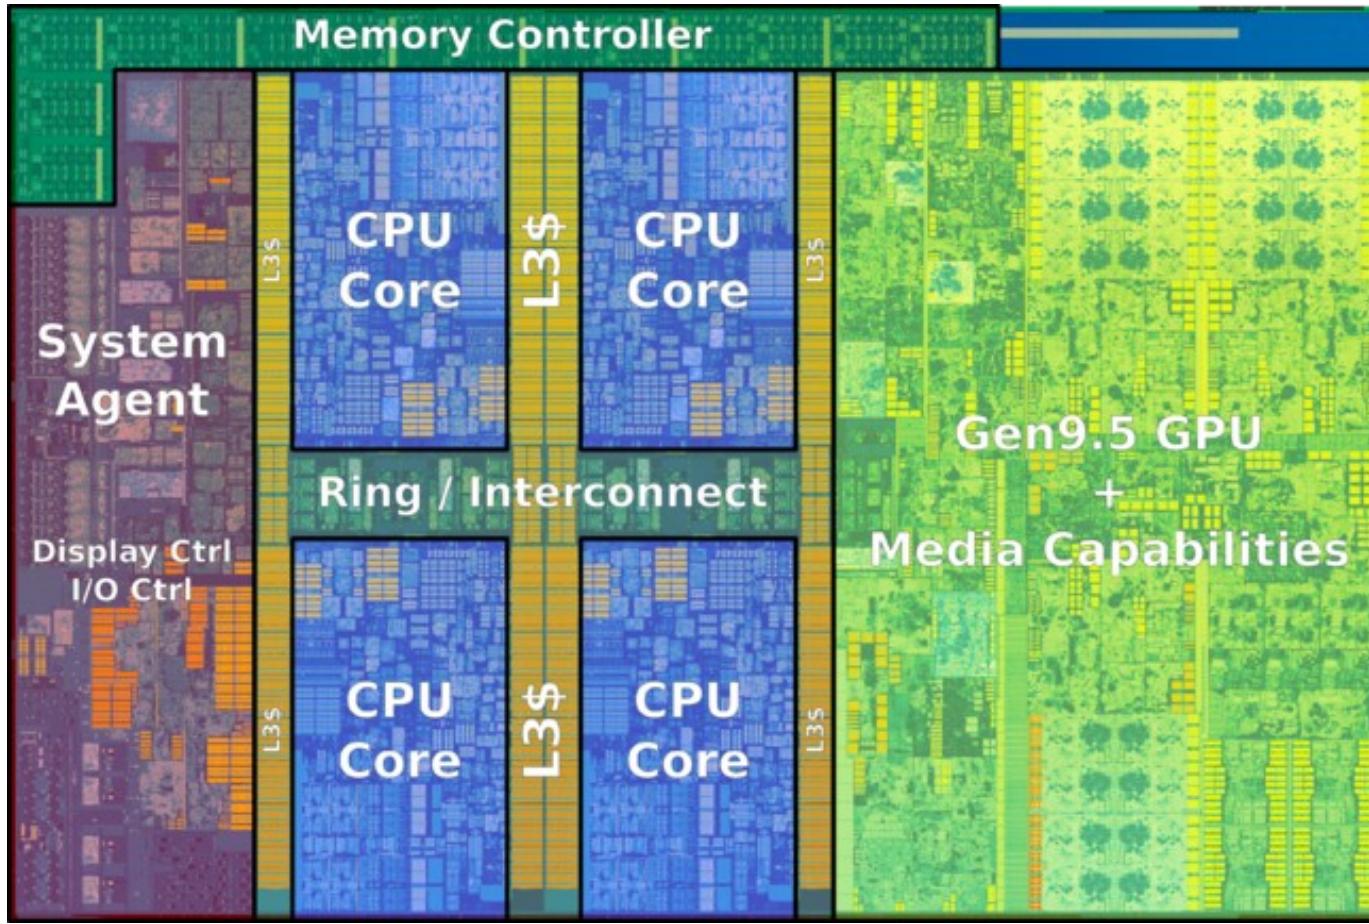


Photo of the quad-core GT2 Kaby Lake processor chip.
Used in mainstream desktop computers.

It has 4 CPU cores to compute in parallel.



Samsung Galaxy S10+

Exynos 9 Series (9820)

6.4" Quad HD+ Curved Dynamic AMOLED (3040x1440)

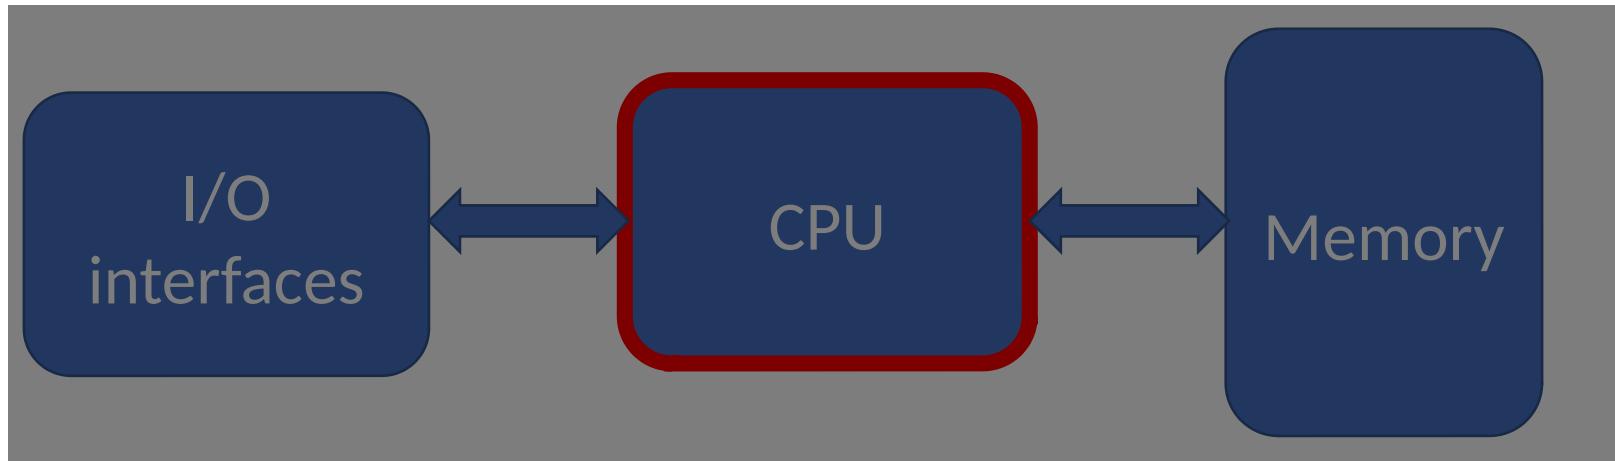
Android 9 (Pie)

8/12GB RAM, 128/256GB/1TB Storage, microSD(up to 512GB)

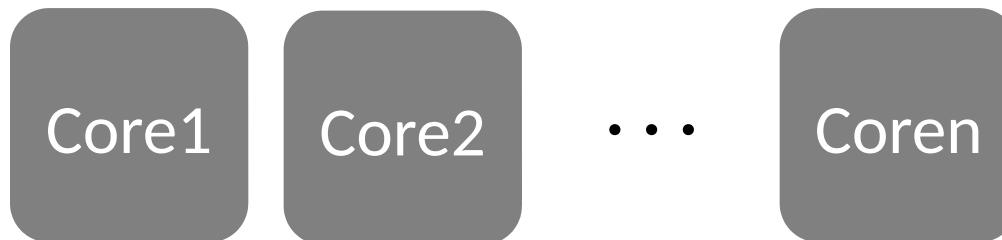
Rear: 12MP AF(Wide)+12MP AF(Tele)+16MP FF(Ultra Wide),
Front: 10MP AF+8MP FF(RGB Depth)

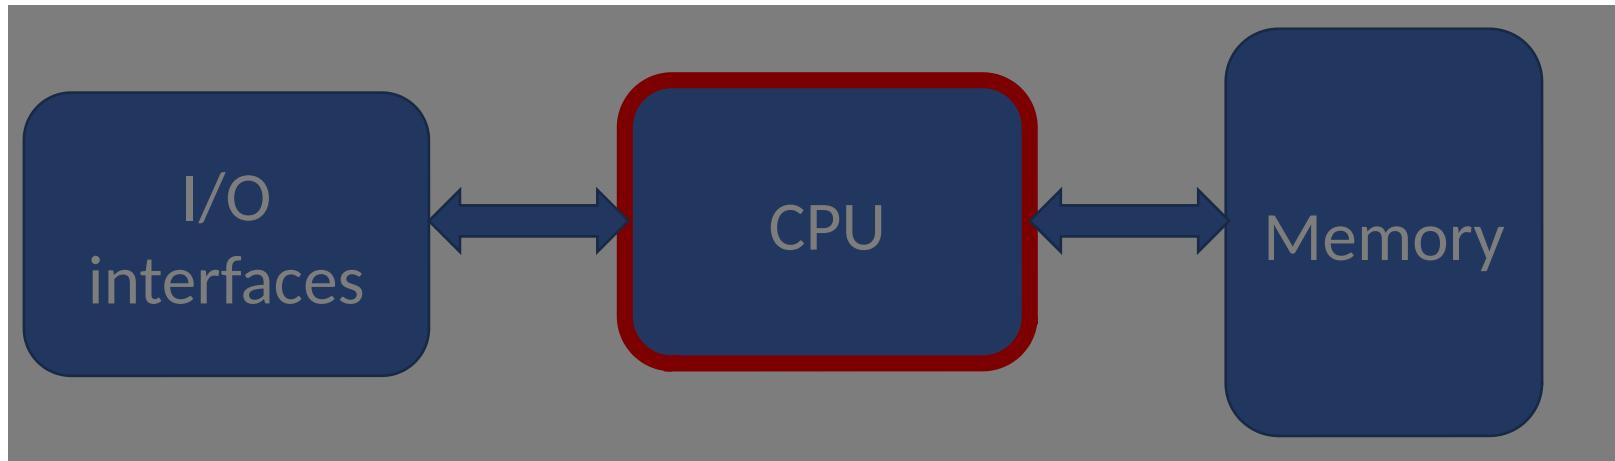
Superior Performance for Seamless Multi-tasking

Featuring a 4th generation custom CPU, the Exynos 9820's innovative tri-cluster architecture delivers premium processing power. The CPU consists of two custom cores for ultimate processing power, two Cortex-A75 cores for optimal performance, and four Cortex-A55 cores for greater efficiency, resulting in superior performance that lasts. Tri-cluster with intelligent task scheduler boosts multi-core performance by 15 percent when compared to the Exynos 9810, while the 4th generation custom CPU with enhanced memory access capability and cutting-edge architecture design improves single core performance by up to 20 percent or boosts power efficiency by up to 40 percent.¹



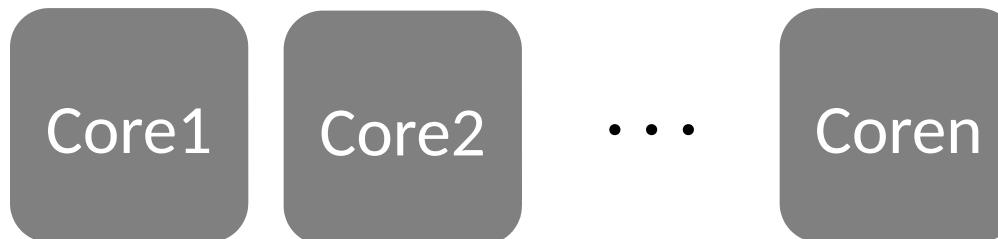
From Single Core
To
Multi Core





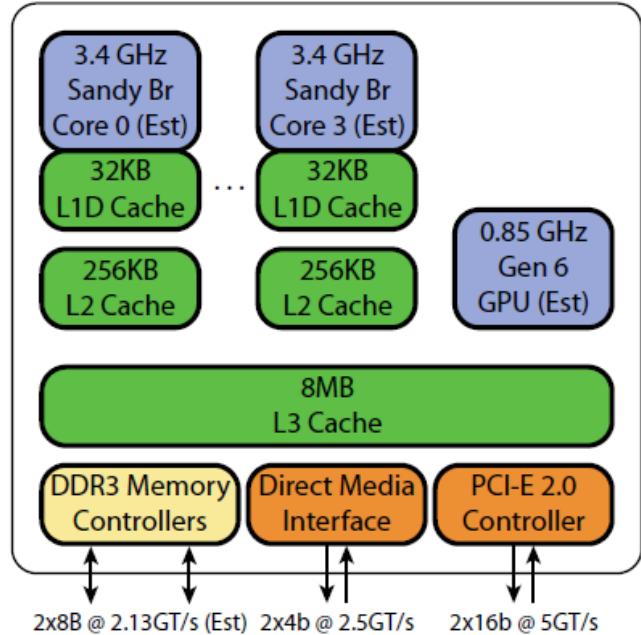
From Single Core
To
Multi Core

What happens to the Memory?
One memory? Multiple memory?

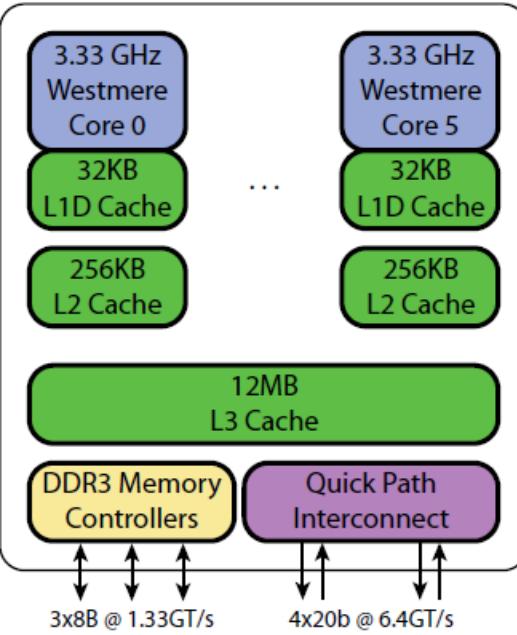


Examples of Intel Multicore Processors

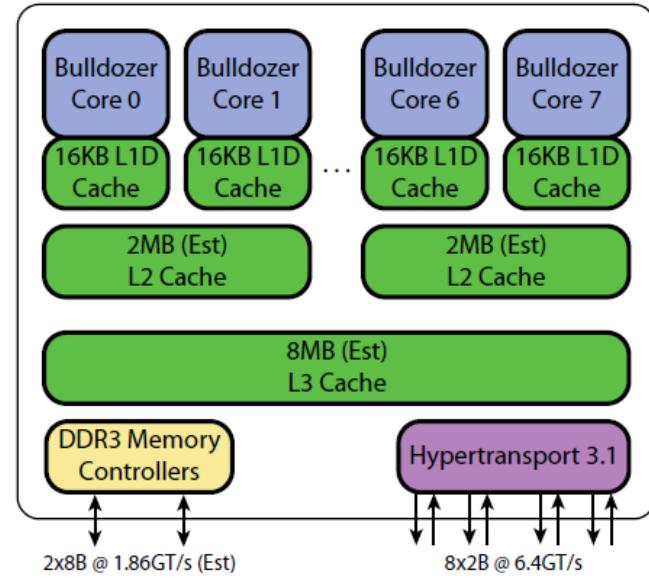
Sandy Bridge Client



Westmere

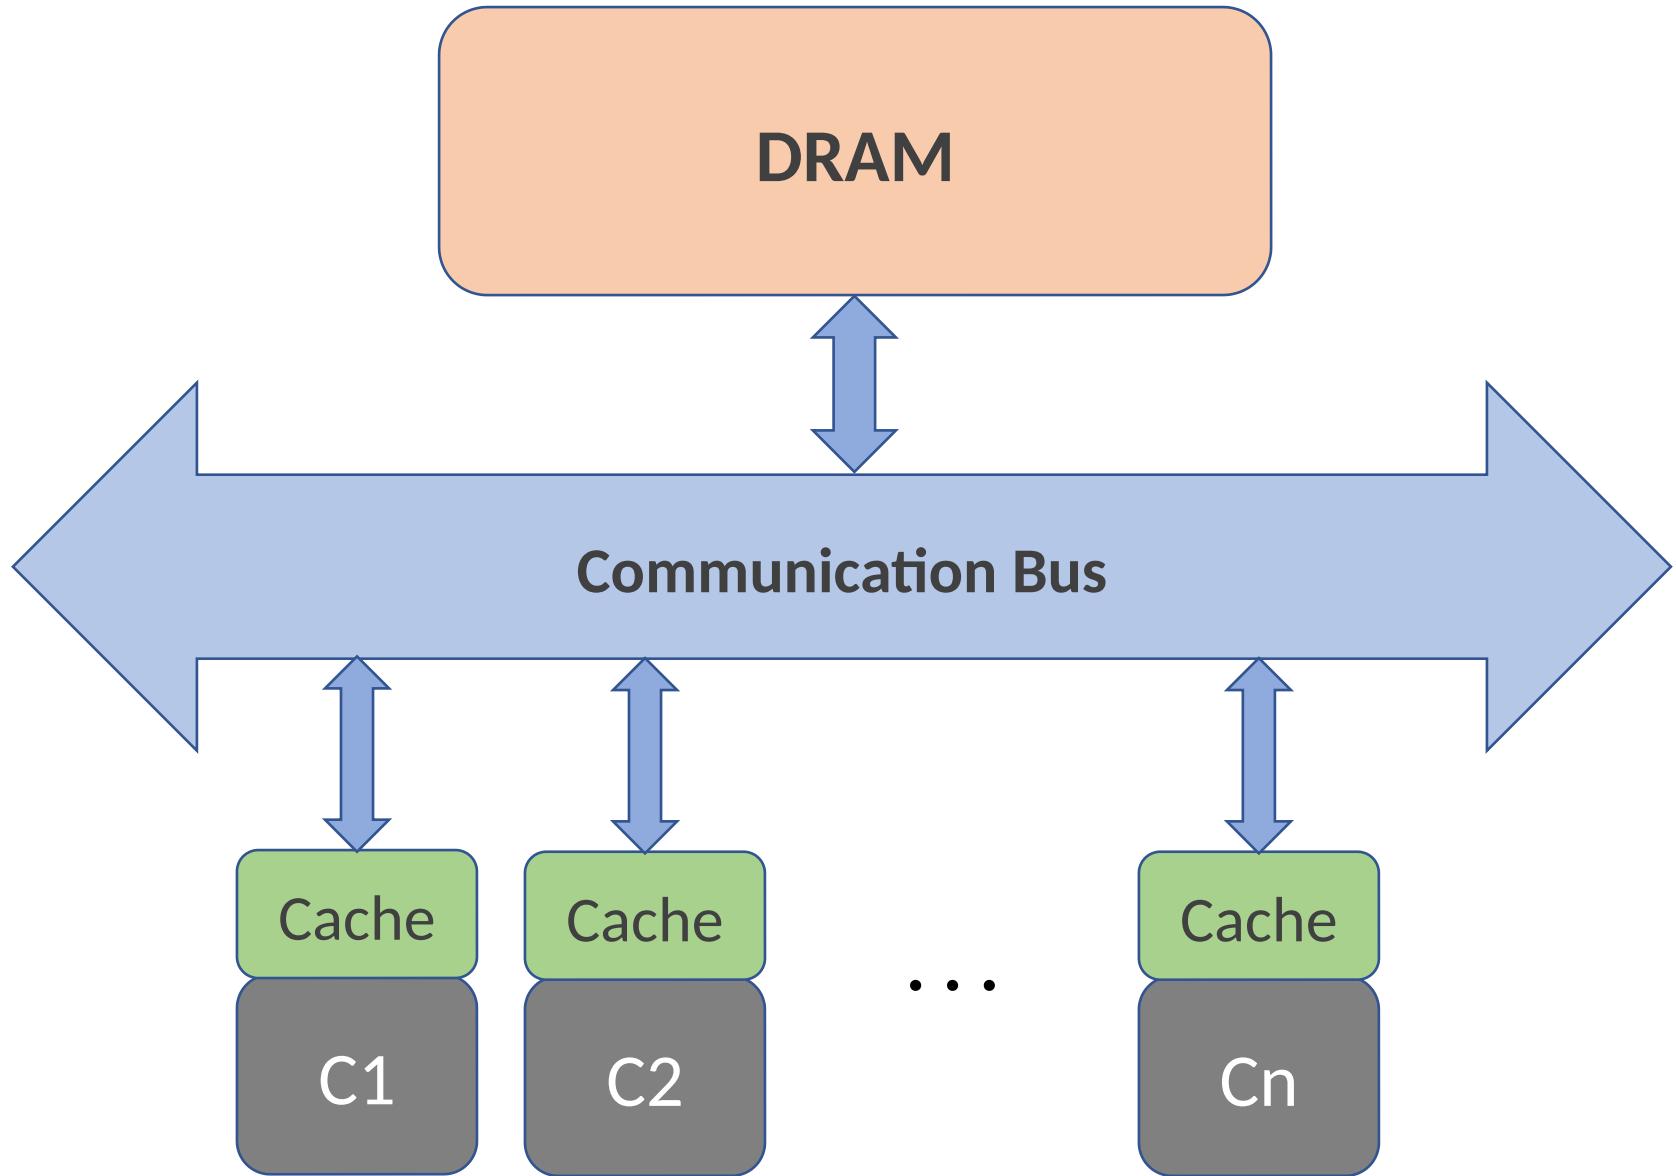


Interlagos Node

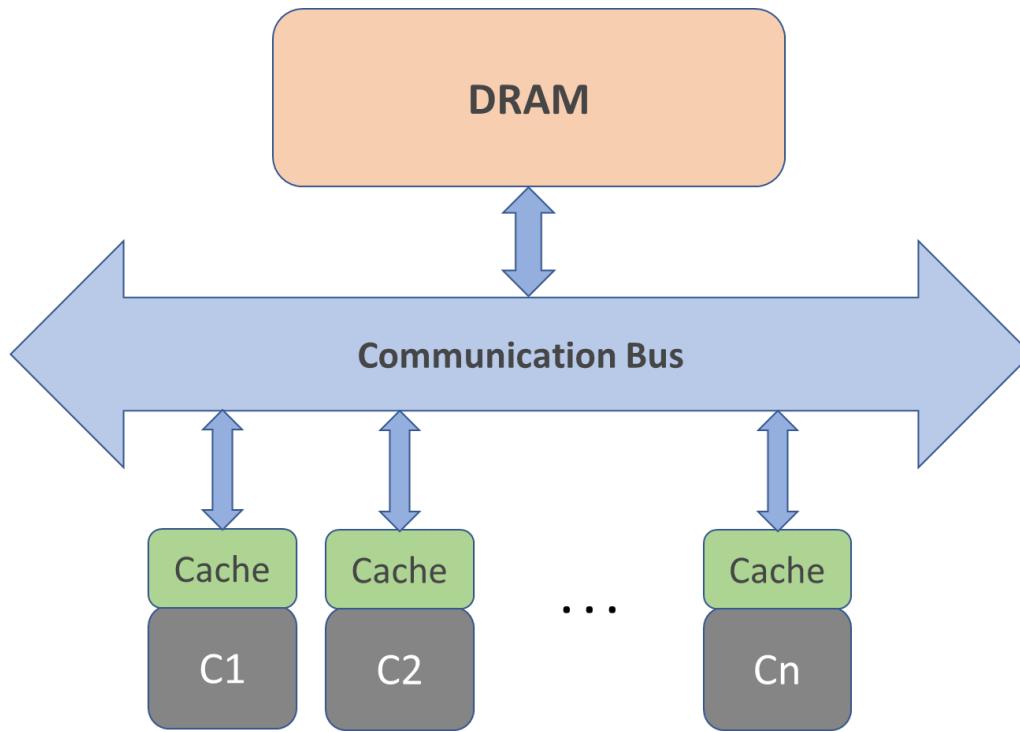


- Each Core has its own L1 Cache
- In some processors, each core has its own L2 Cache
- All cores share one L3 cache
- All cores share one DDR3 DRAM (main memory)

Simplified Multicore Architecture



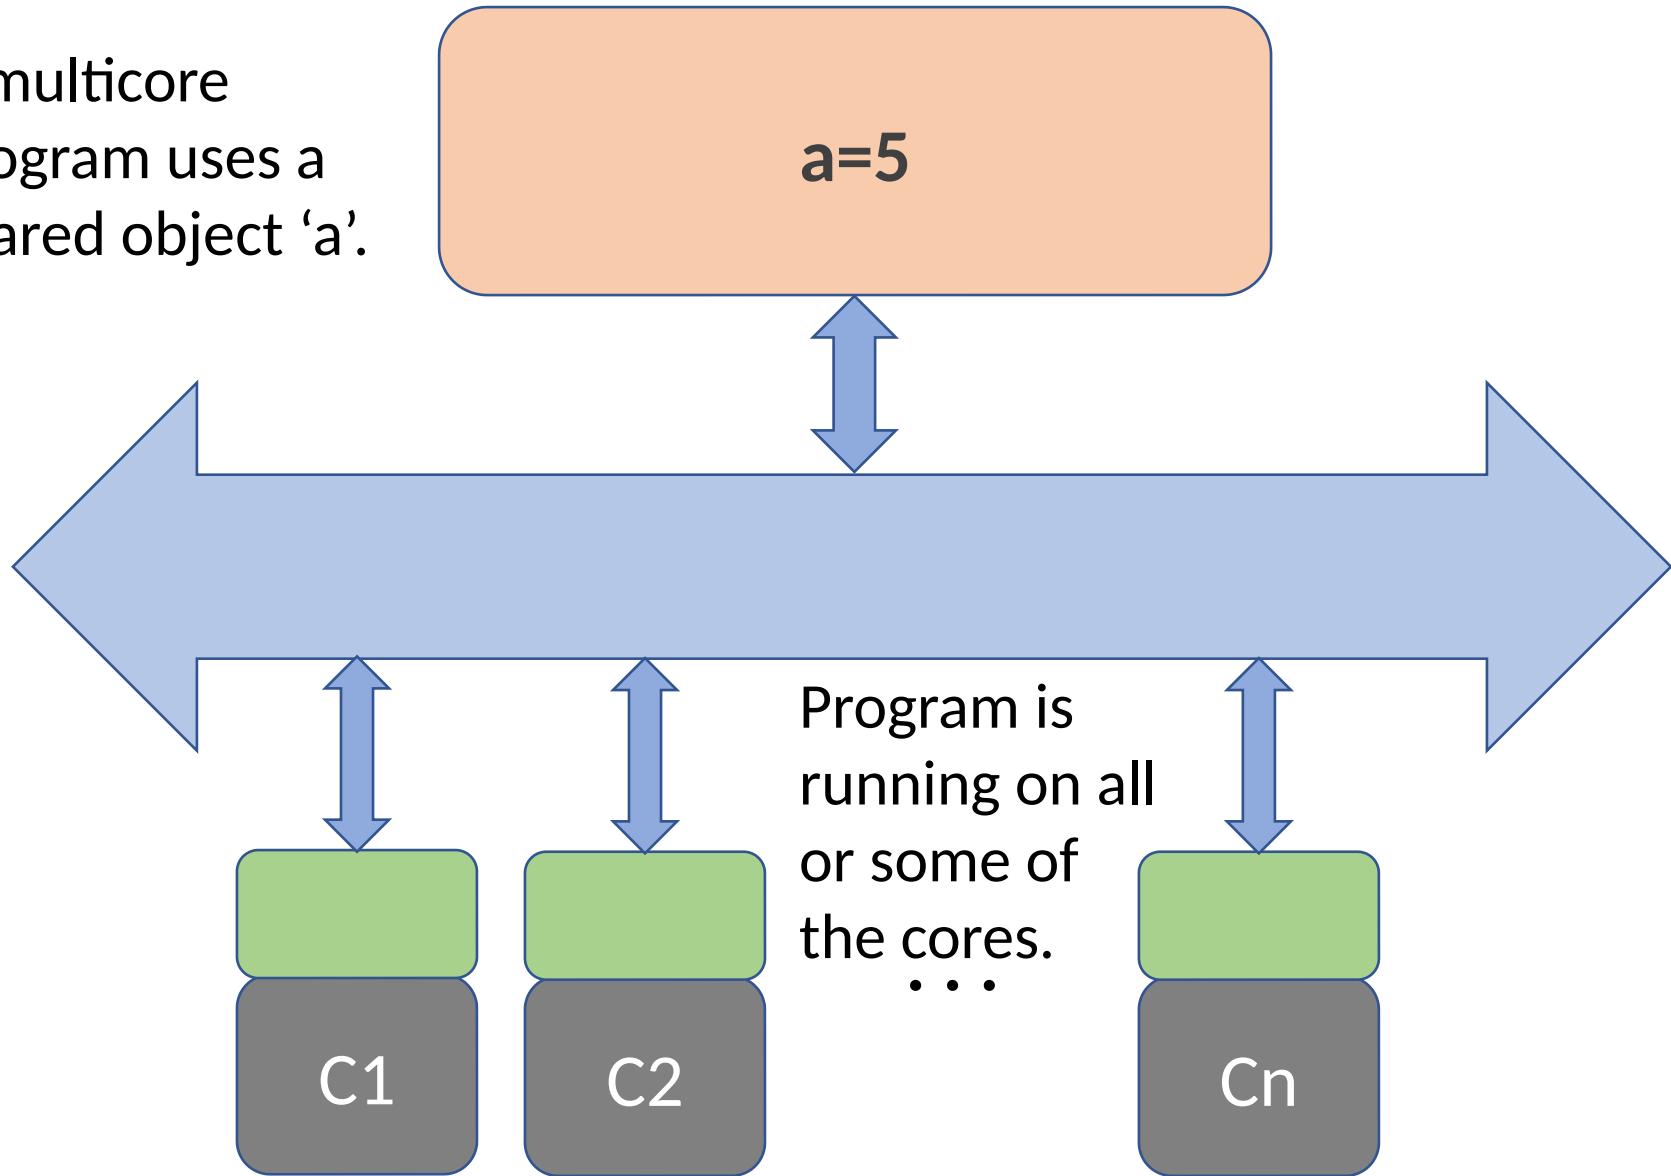
Coherence of data



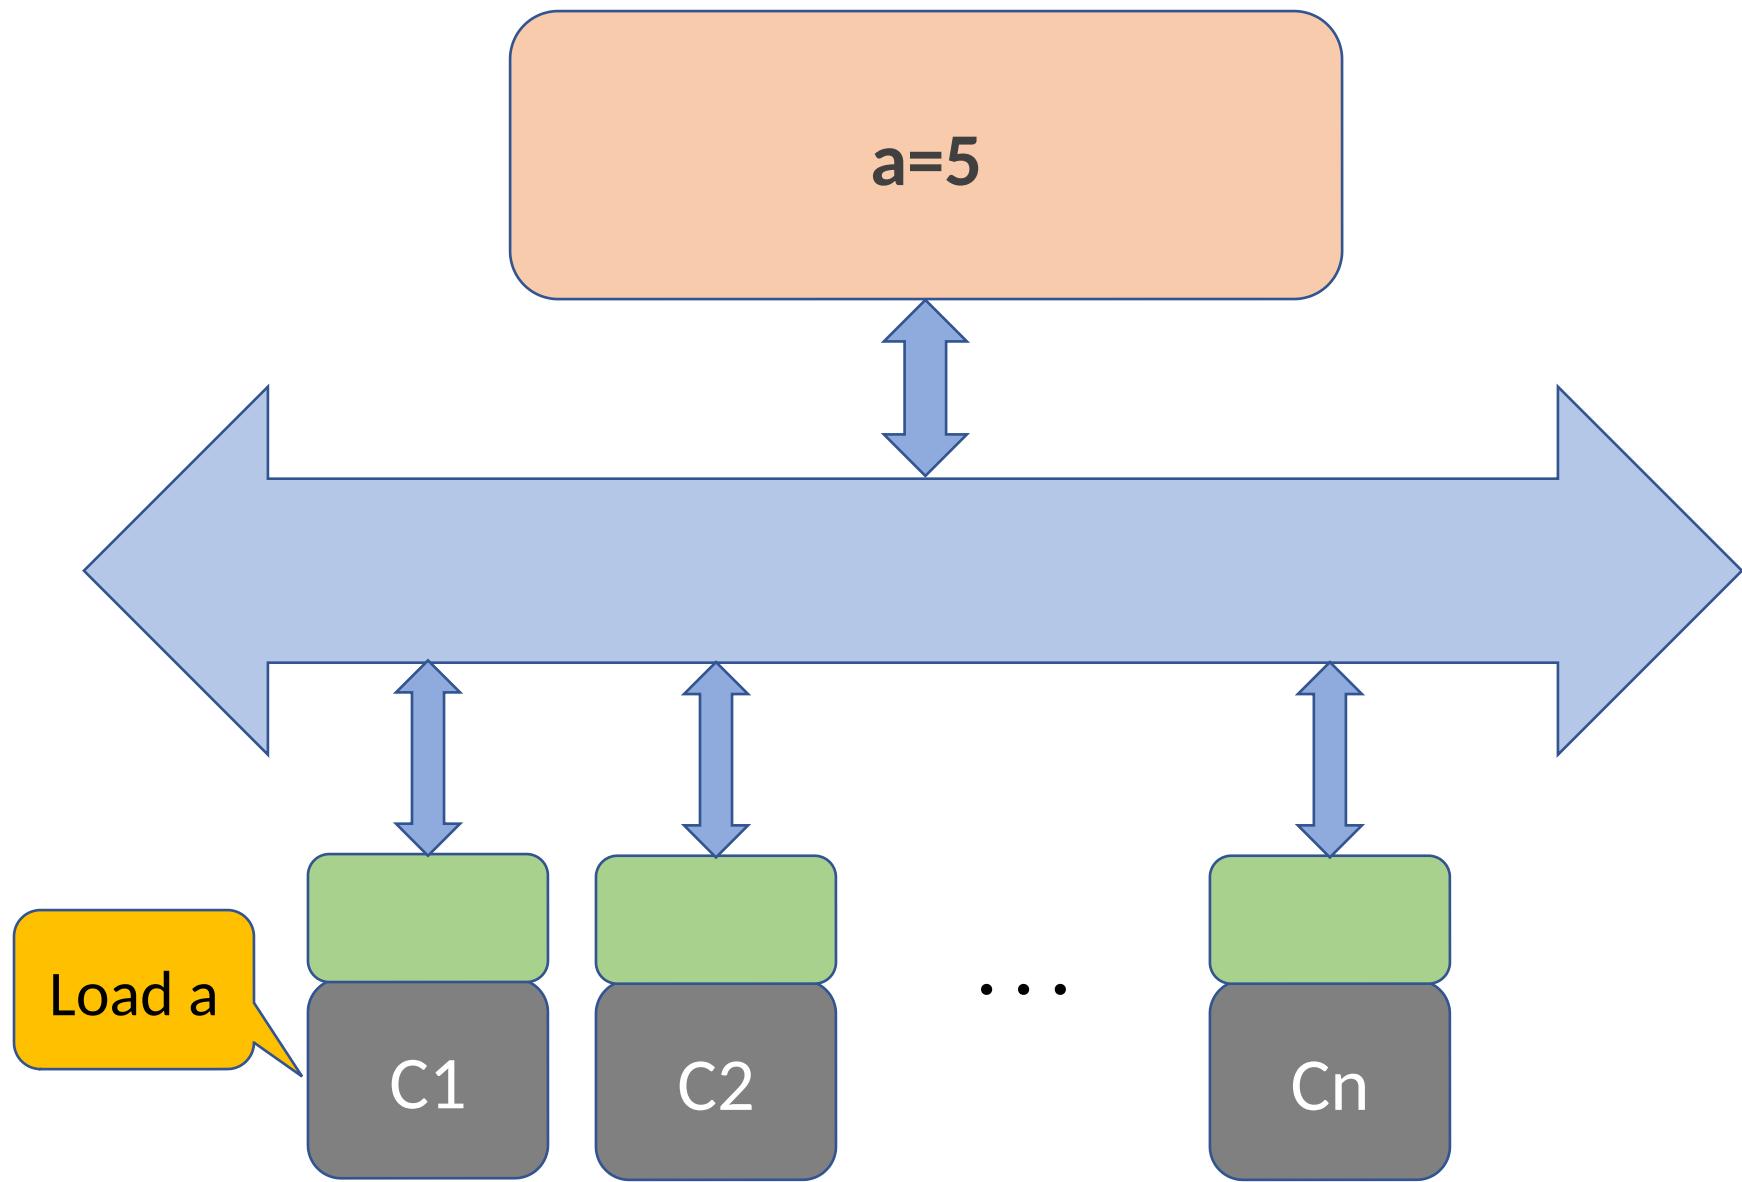
- ‘Coherence’ is the quality of being logical and consistent. During program execution, data must remain coherent.
- On a multicore system, a data variable may reside in multiple caches and might get updated ‘locally’ by its local core
✉ this causes coherence problem

Example: Cache coherence problem

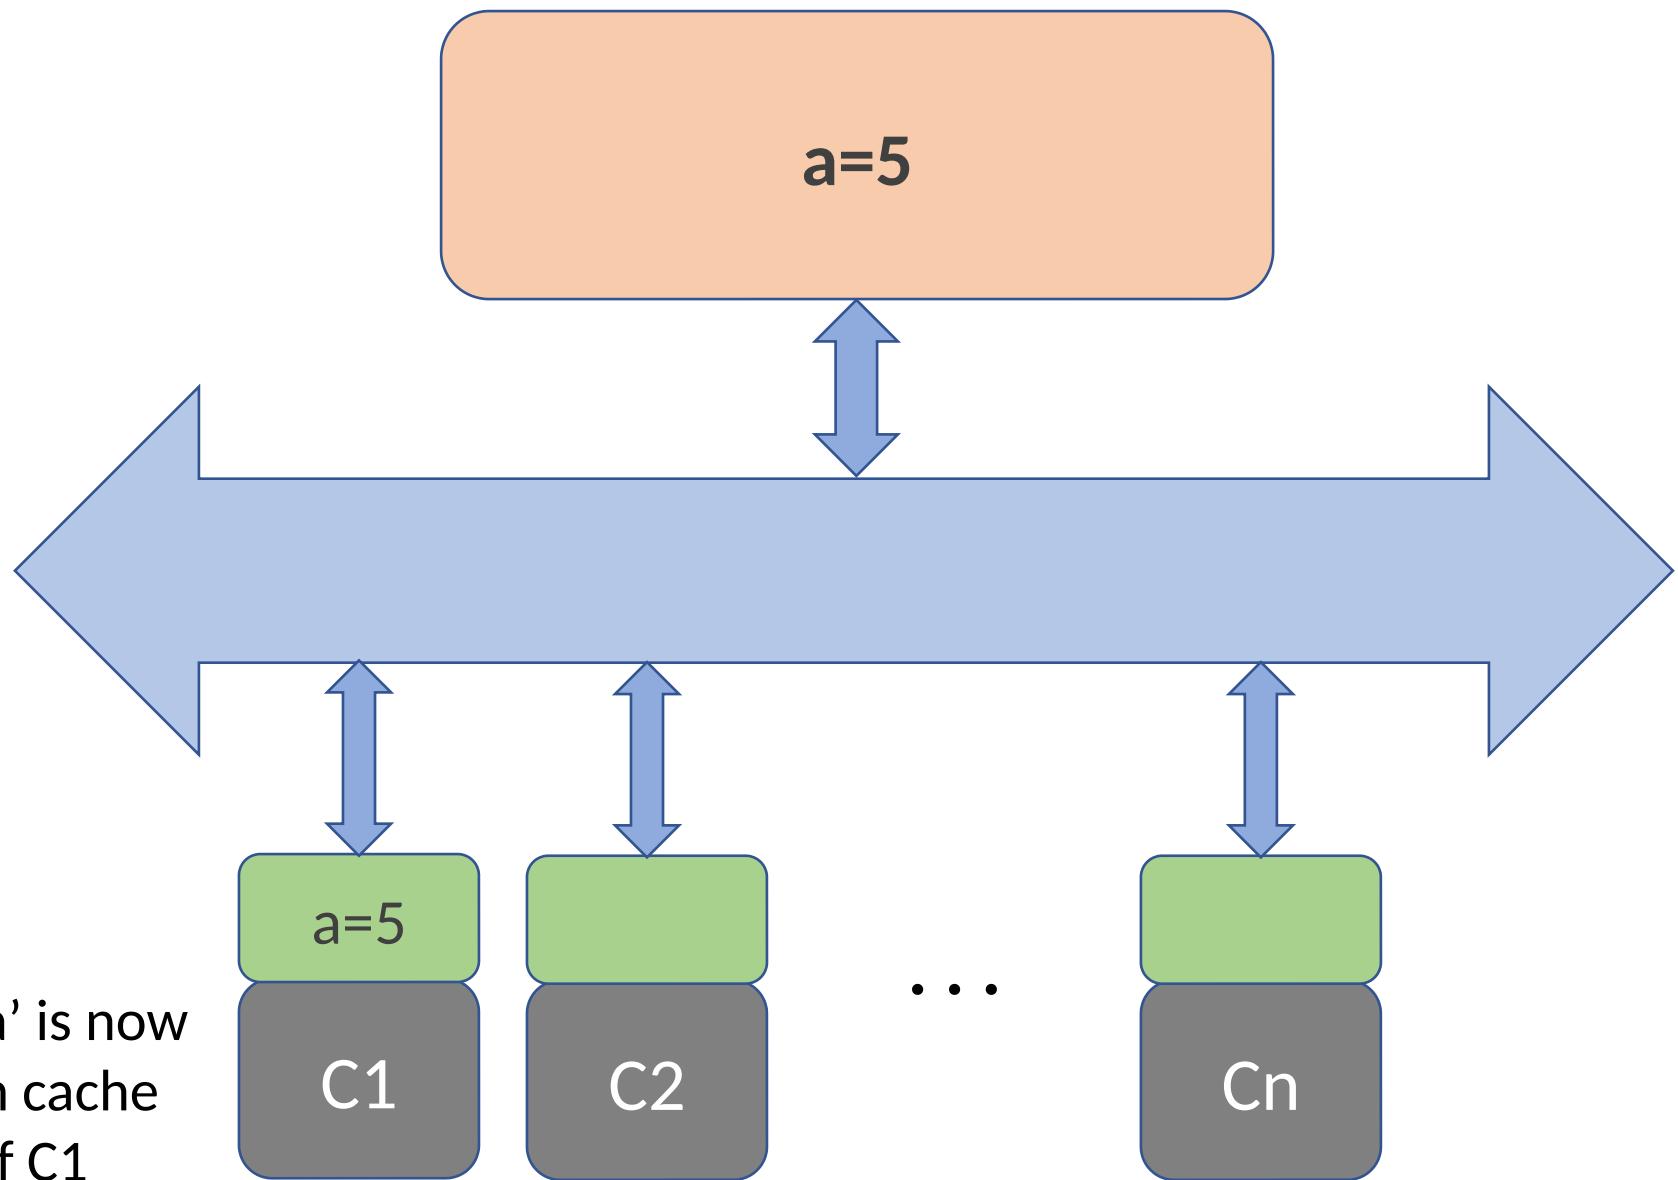
A multicore program uses a shared object 'a'.



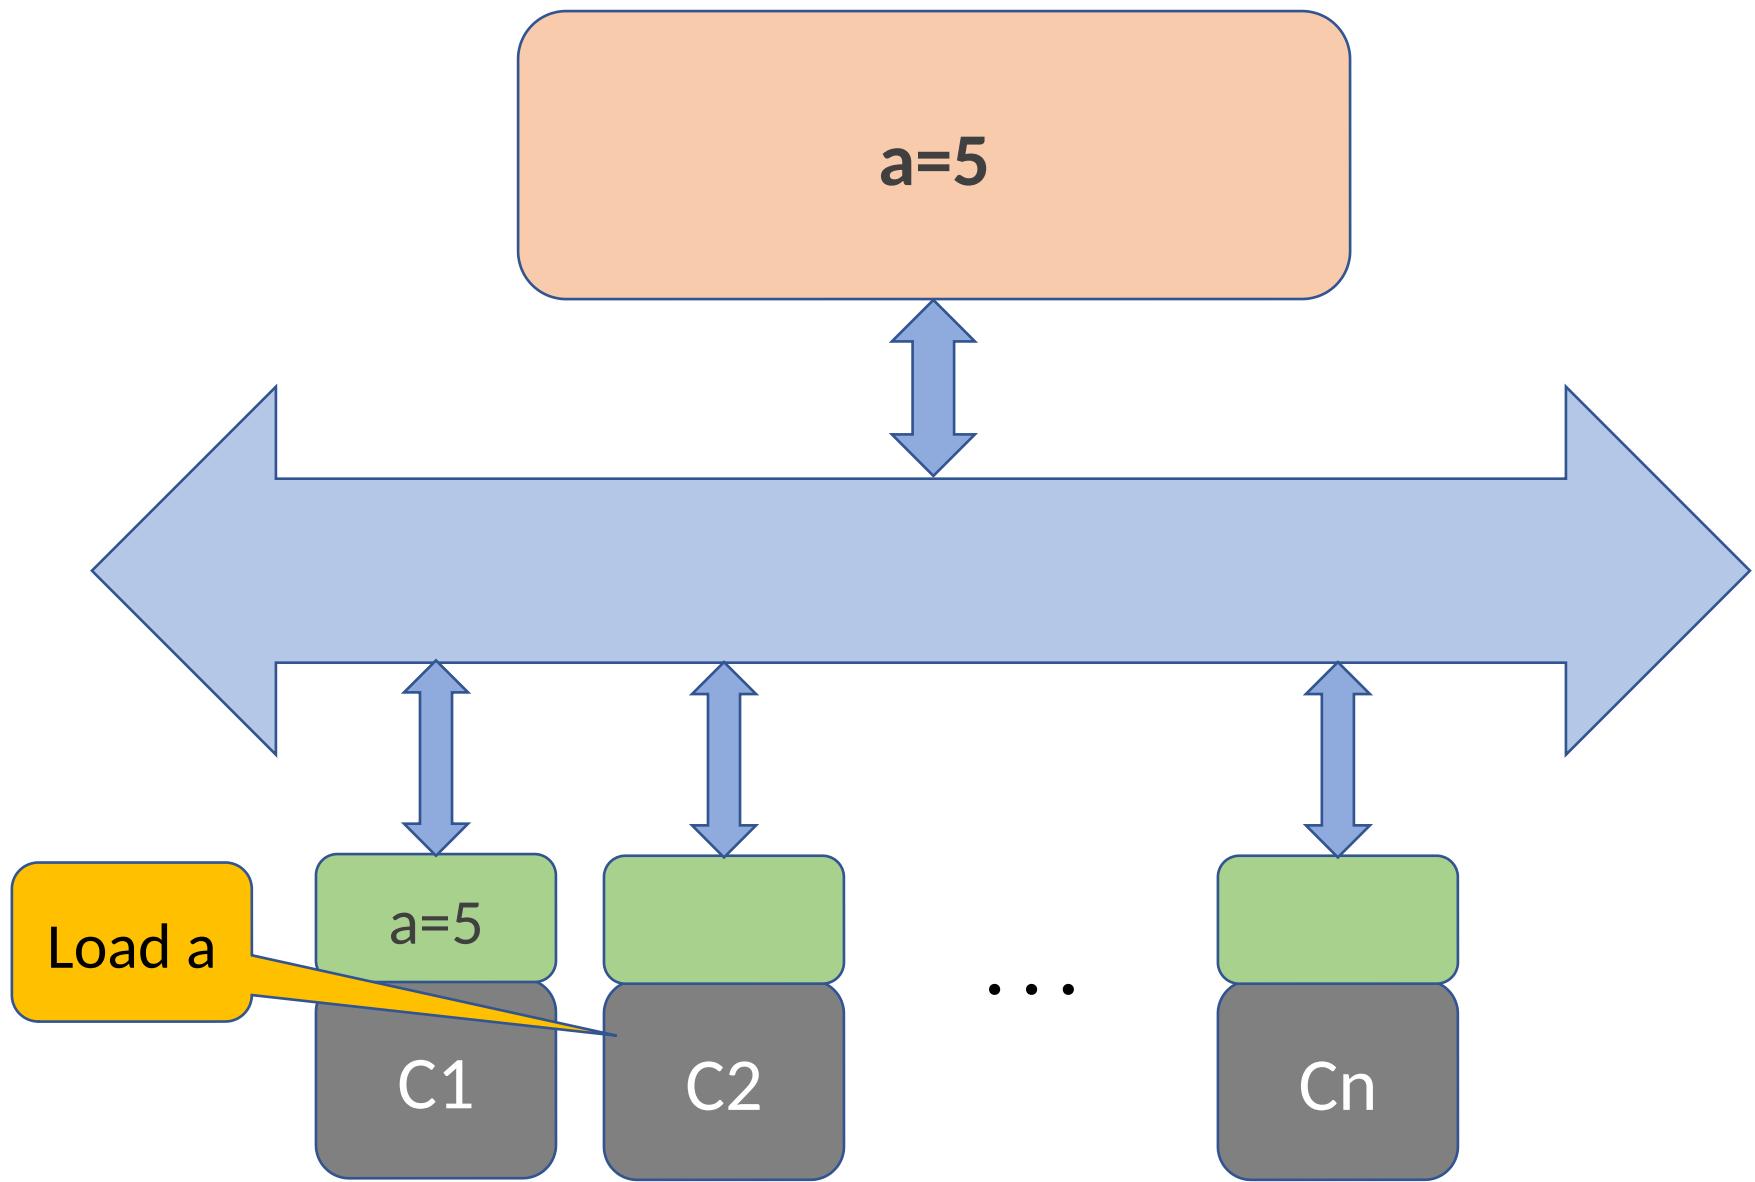
Example: Cache coherence problem



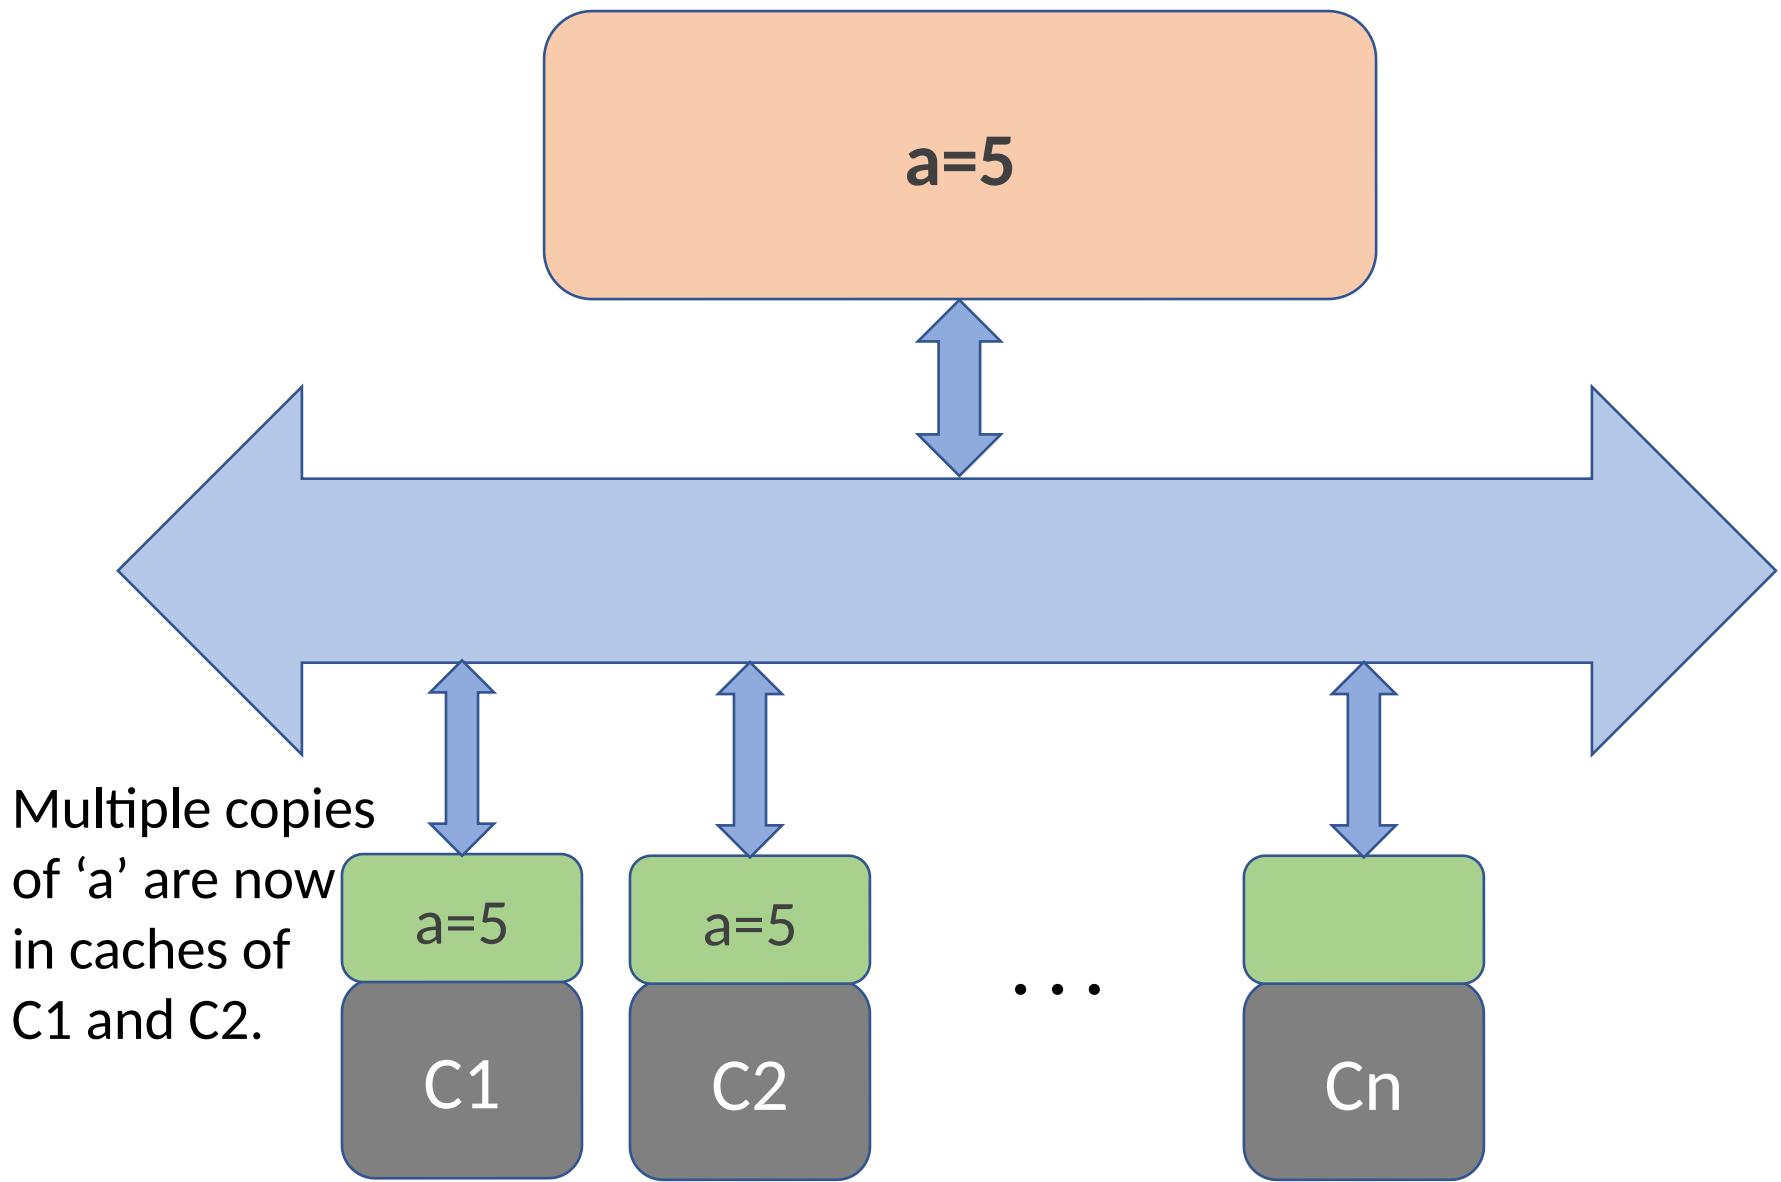
Example: Cache coherence problem



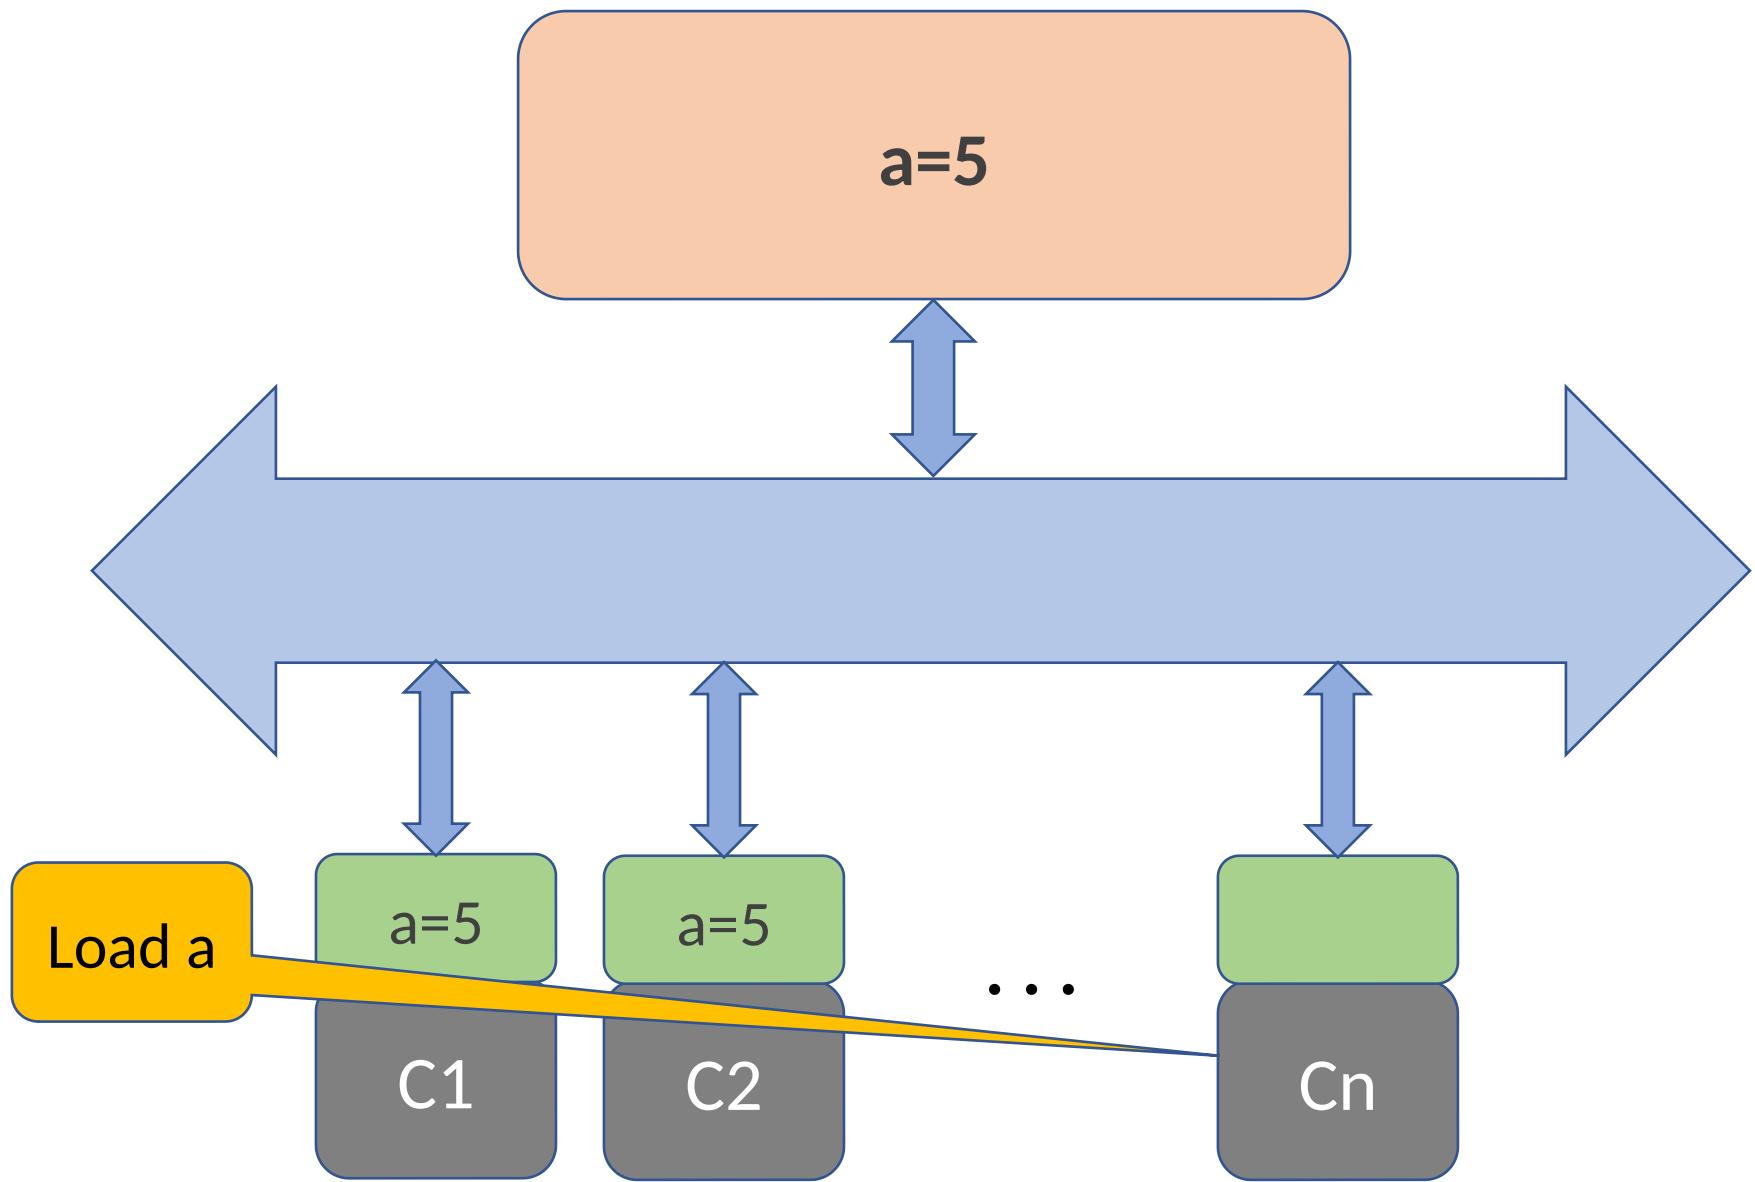
Example: Cache coherence problem



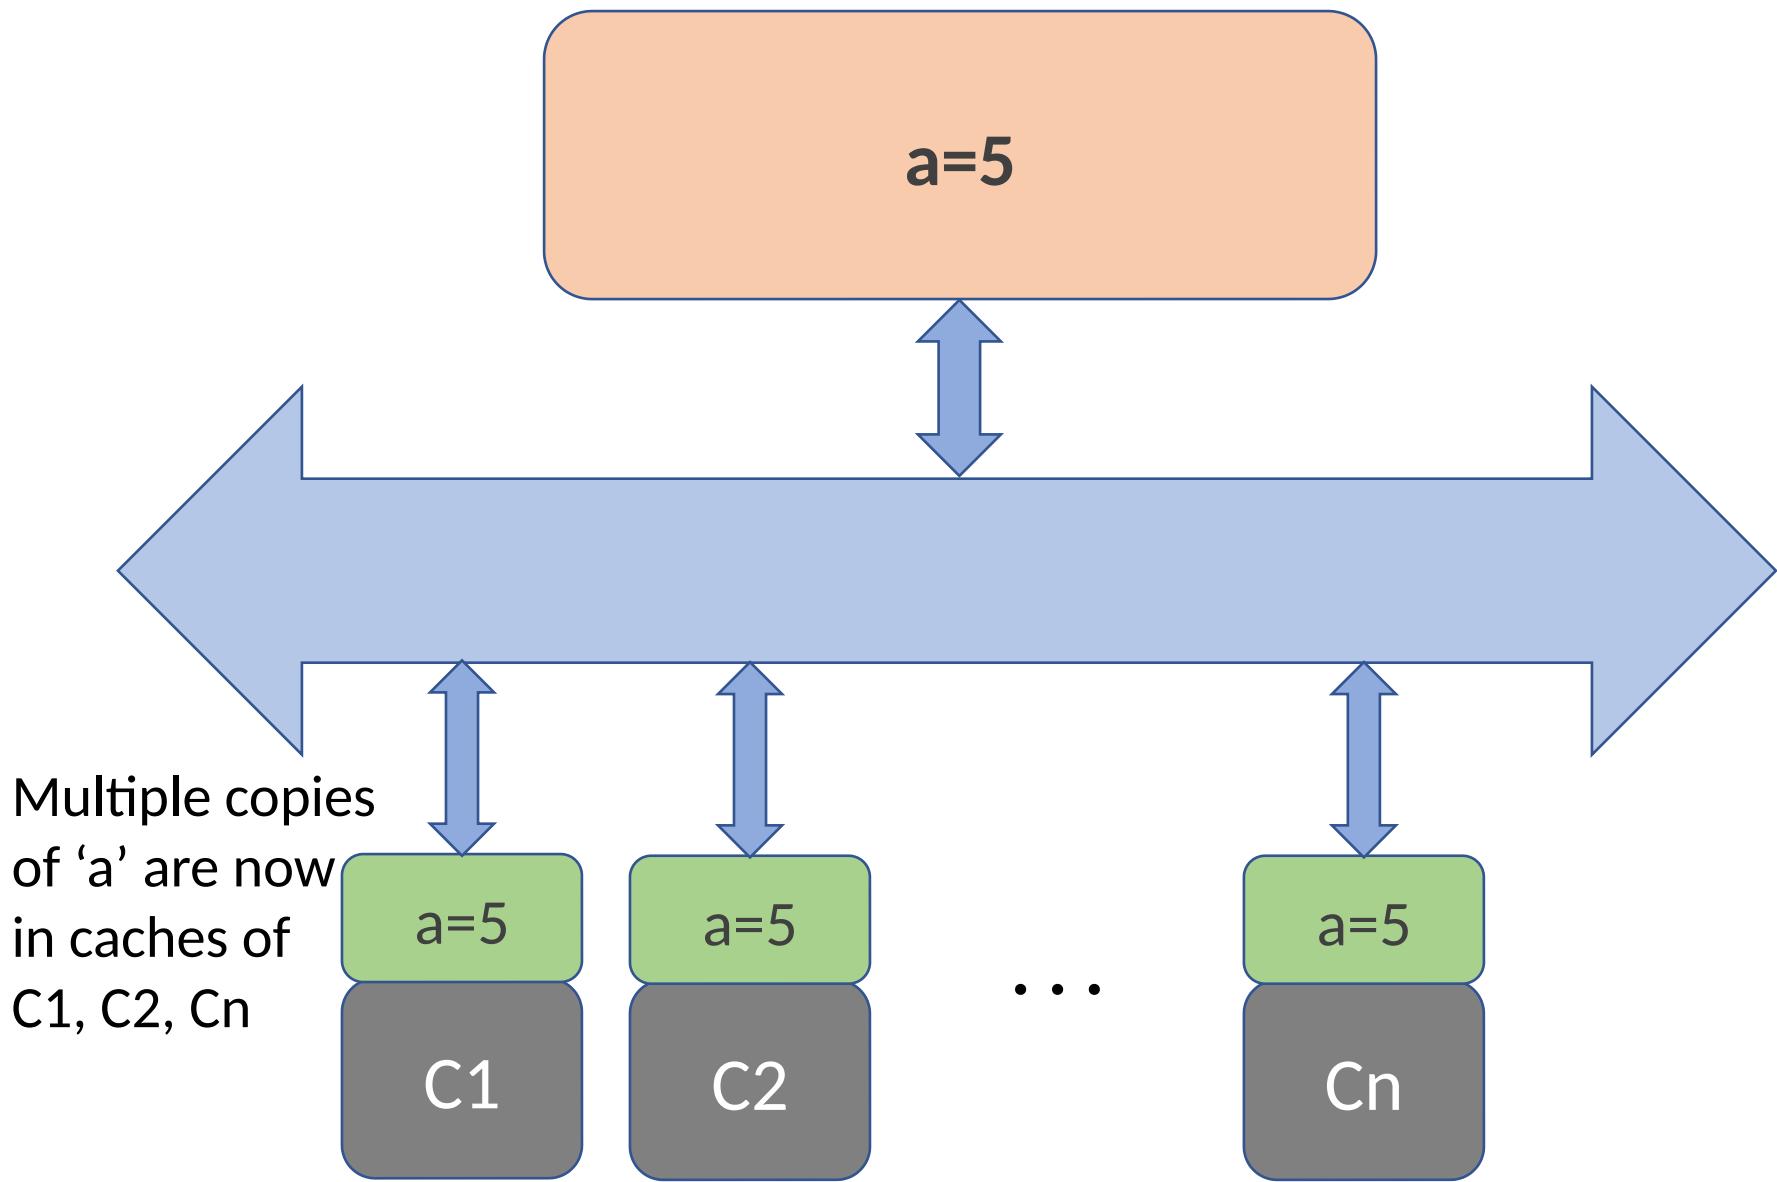
Example: Cache coherence problem



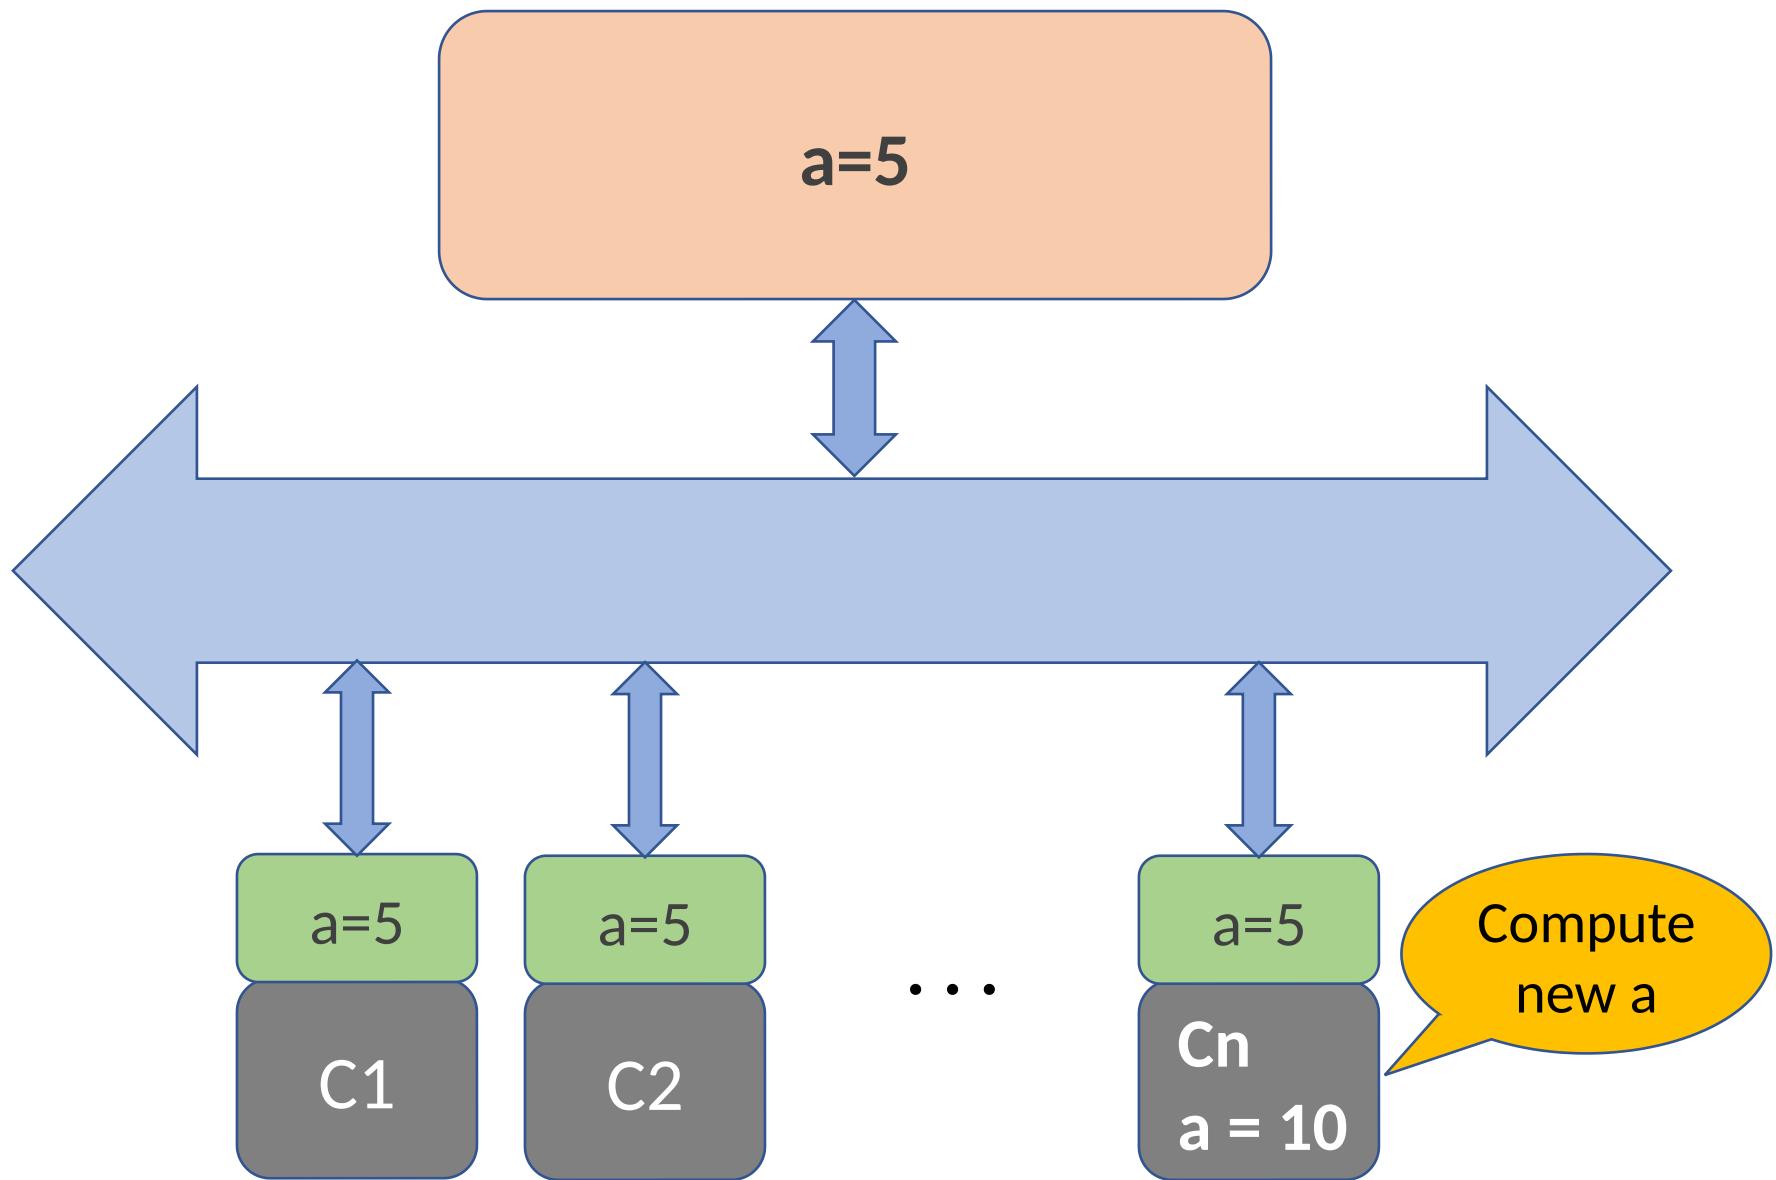
Example: Cache coherence problem



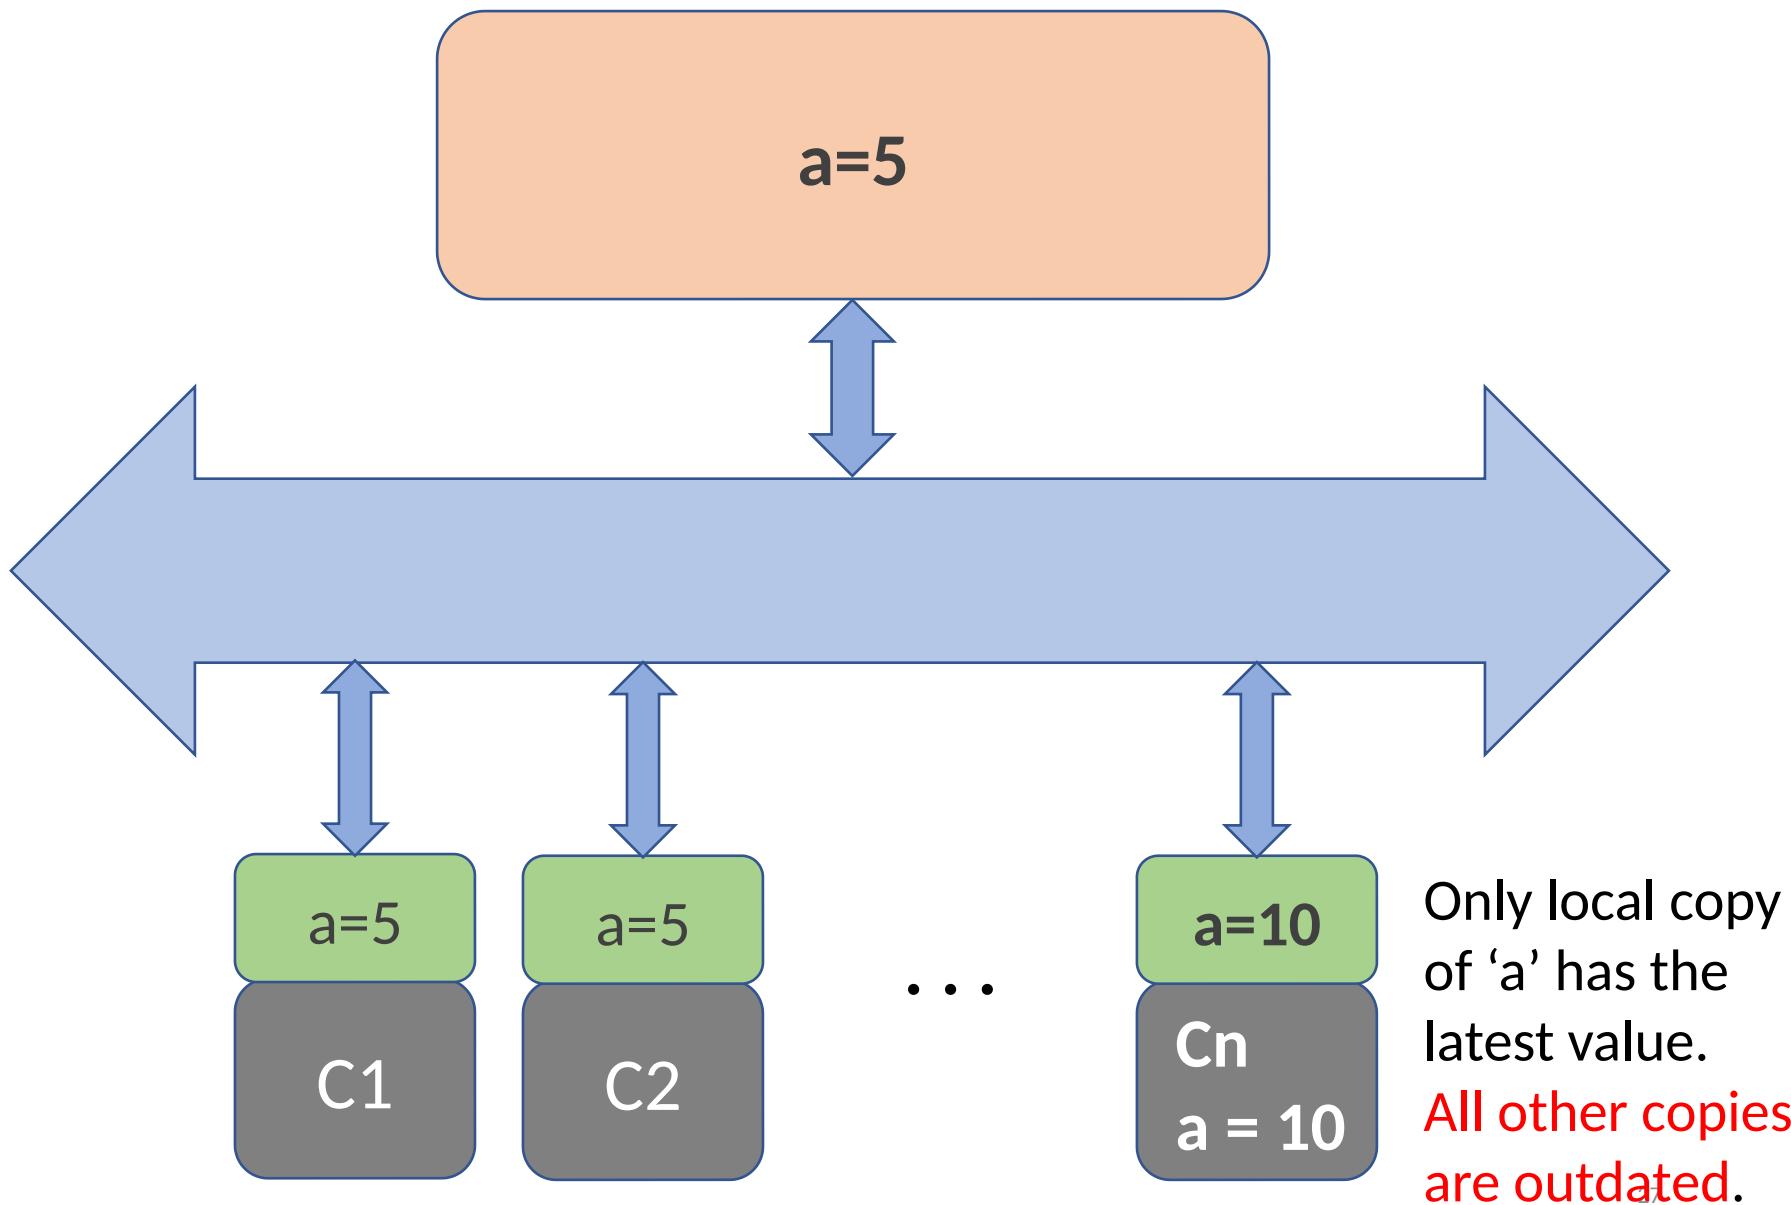
Example: Cache coherence problem



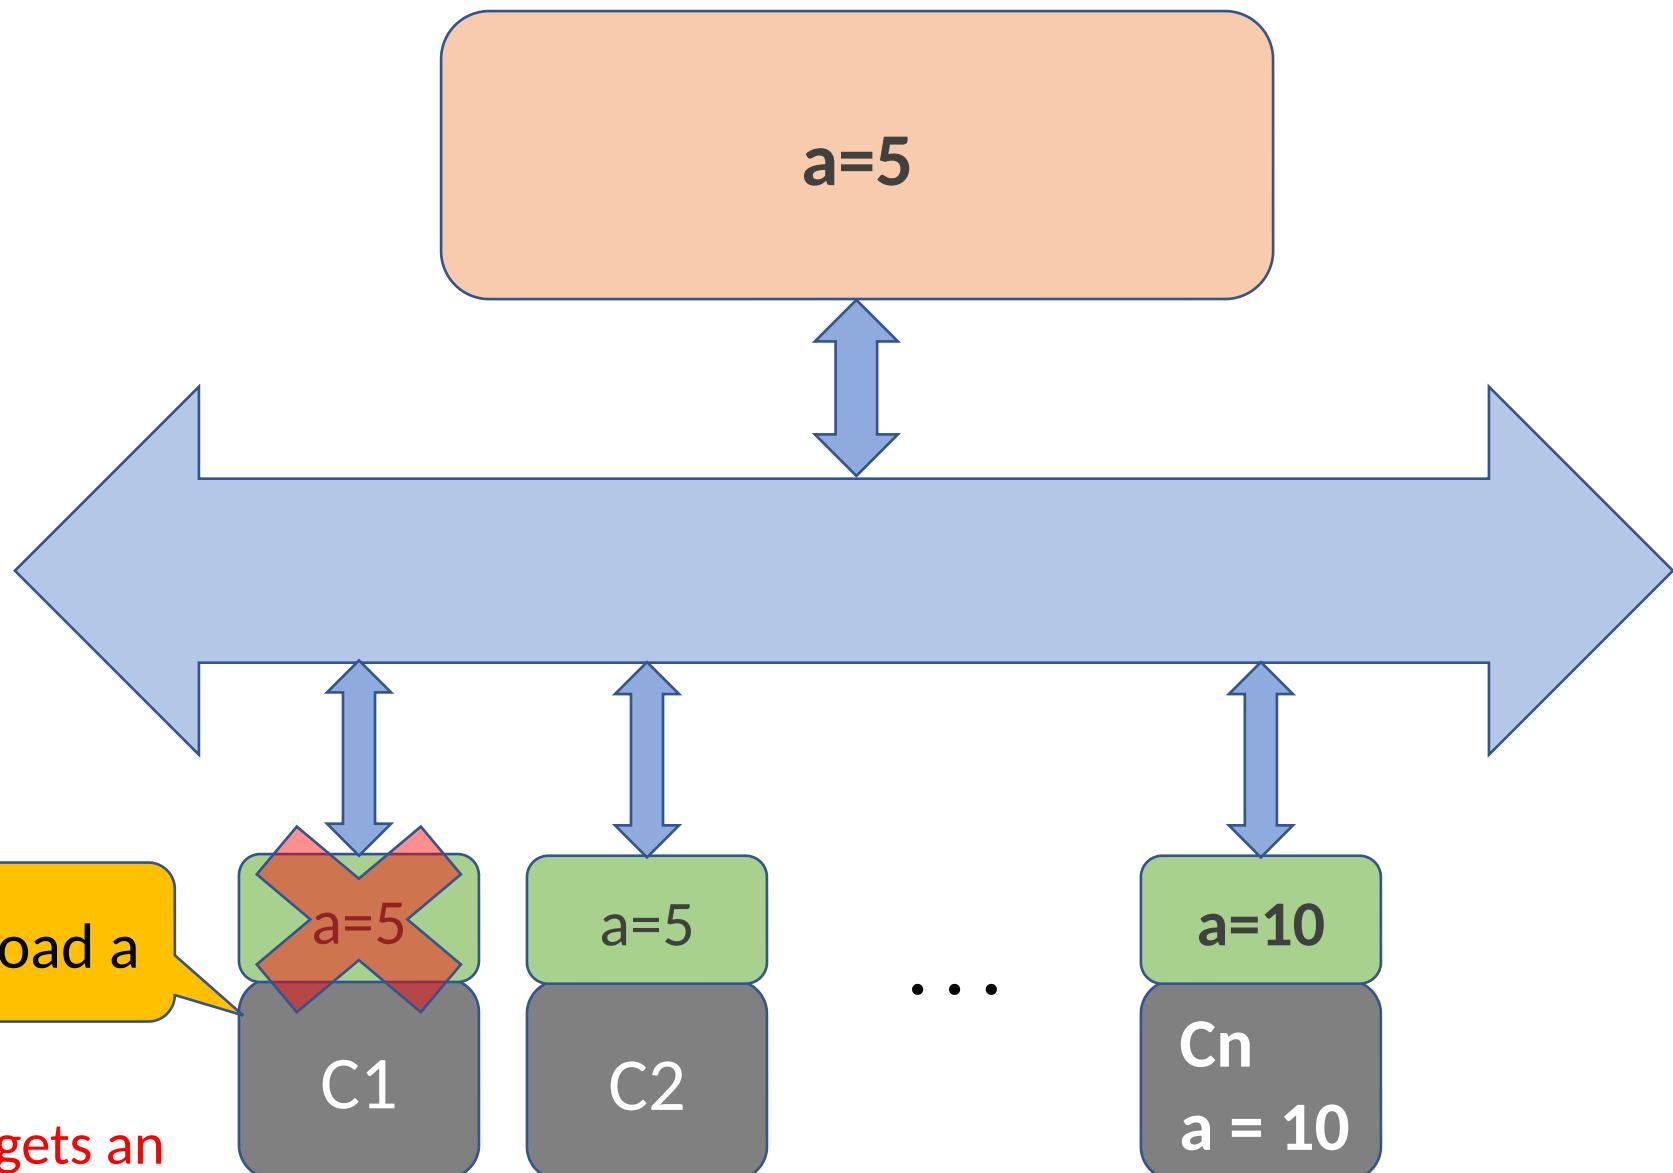
Example: Cache coherence problem



Example: Cache coherence problem



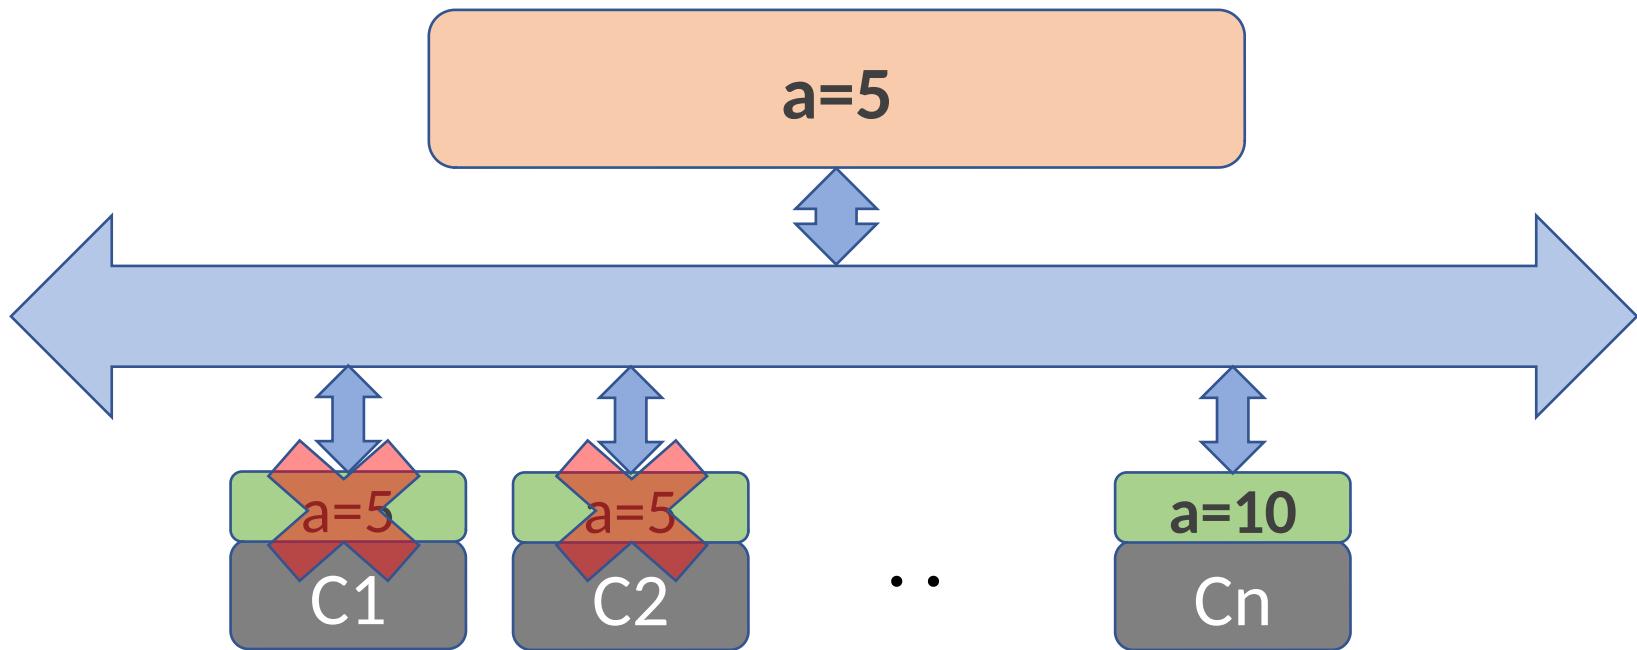
Example: Cache coherence problem



Coherence mechanisms

To achieve coherence:

a **write to a memory location** must cause all other copies of this location to be **removed from the caches** they are in.



Do: As soon as 'a' gets updated by a core, invalidate or delete all other copies of it.

Coherence protocols

- Coherence protocols apply cache coherence in multicore systems.
- Goal is that two cores **must never see different values** for the same shared data.

How this is achieved?

Additional information is stored and passed

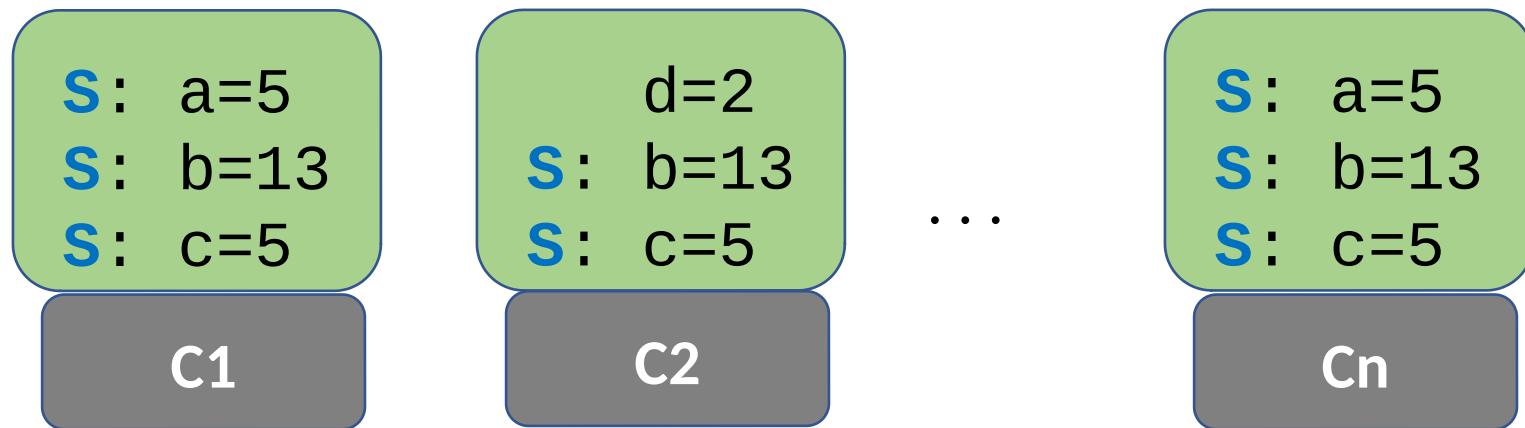
- Is this data stored in multiple caches?
- Has this cache line been modified?

MSI protocols

MSI protocol is a simple cache coherence protocol.
In this protocol, each cache line is labeled with a state:

- **M**: cache block has been modified.
- **S**: other caches may be sharing this block.
- **I**: cache block is invalid

Initial state of caches before computation

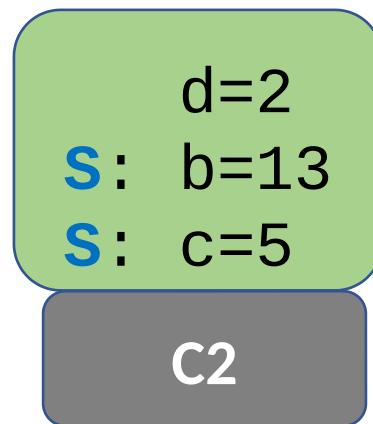
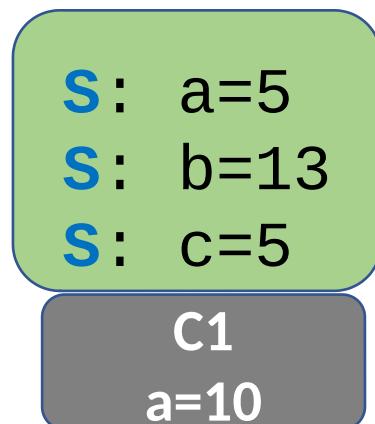


MSI protocols

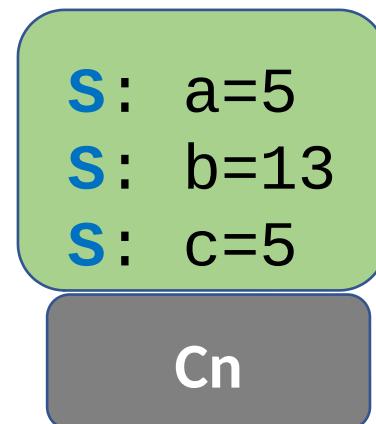
MSI protocol is a simple cache coherence protocol.
In this protocol, each cache line is labeled with a state:

- **M**: cache block has been modified.
- **S**: other caches may be sharing this block.
- **I**: cache block is invalid

C1 computes new value a=10 in register



...

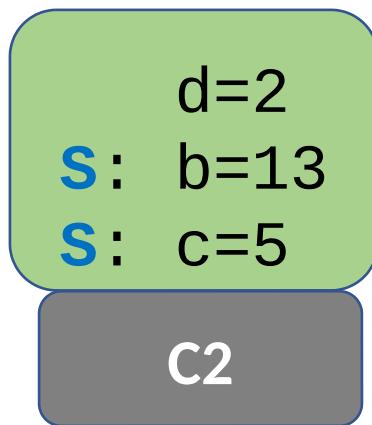
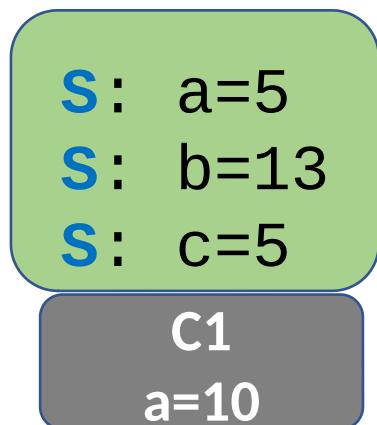


MSI protocols

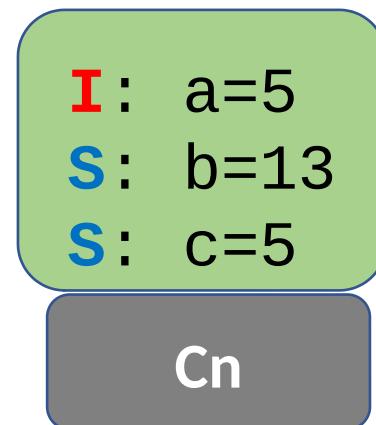
MSI protocol is a simple cache coherence protocol.
In this protocol, each cache line is labeled with a state:

- **M**: cache block has been modified.
- **S**: other caches may be sharing this block.
- **I**: cache block is invalid

Hardware invalidates all other copies of a



...

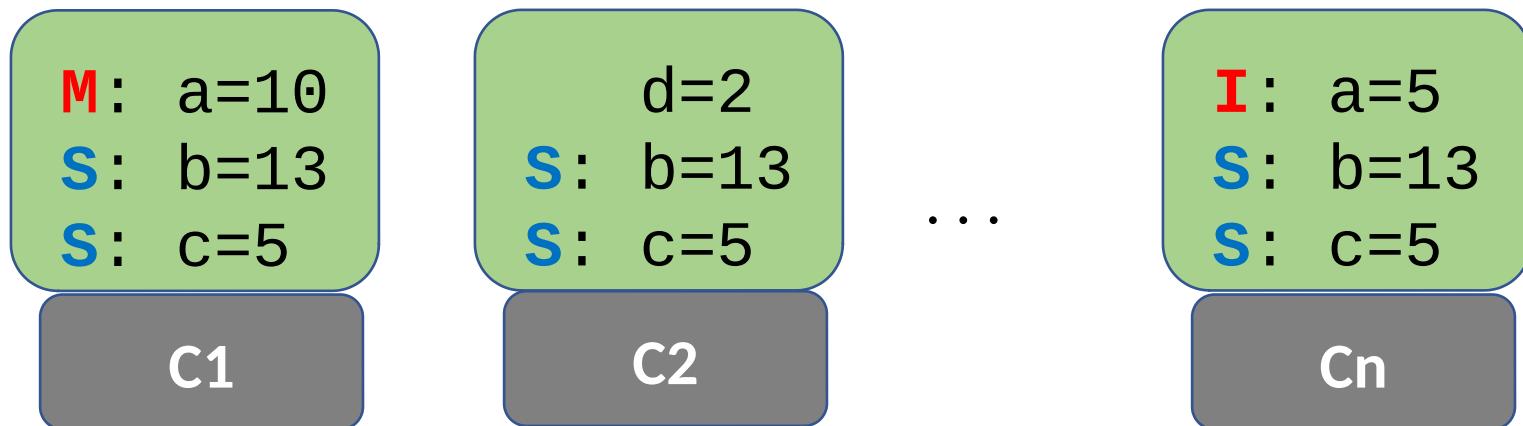


MSI protocols

MSI protocol is a simple cache coherence protocol.
In this protocol, each cache line is labeled with a state:

- **M**: cache block has been modified.
- **S**: other caches may be sharing this block.
- **I**: cache block is invalid

Hardware updates a in cache of C1 and marks it with 'M'

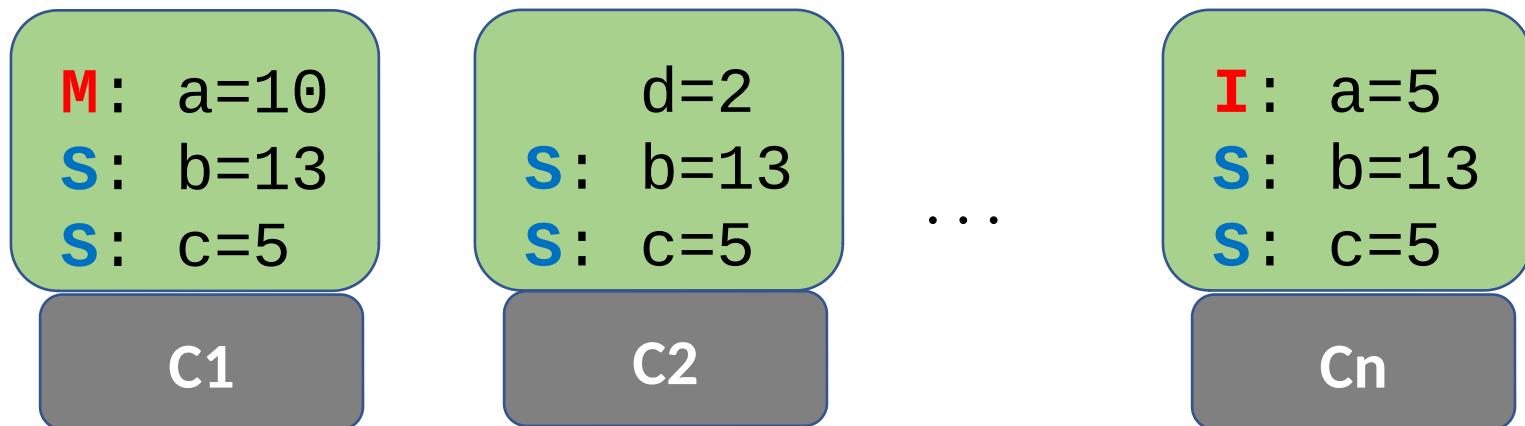


MSI protocols

MSI protocol is a simple cache coherence protocol.
In this protocol, each cache line is labeled with a state:

- **M**: cache block has been modified.
- **S**: other caches may be sharing this block.
- **I**: cache block is invalid

Hardware also updates 'a' in the shared memory

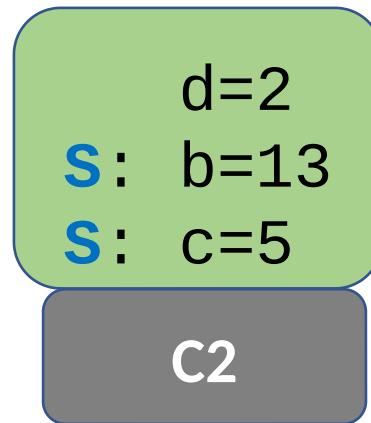
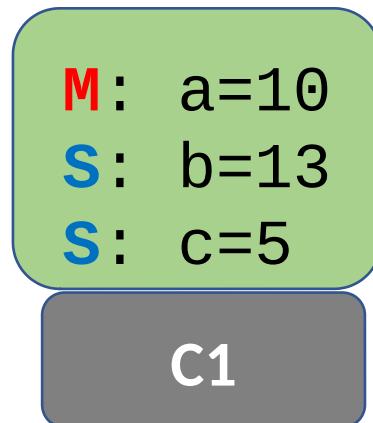


MSI protocols

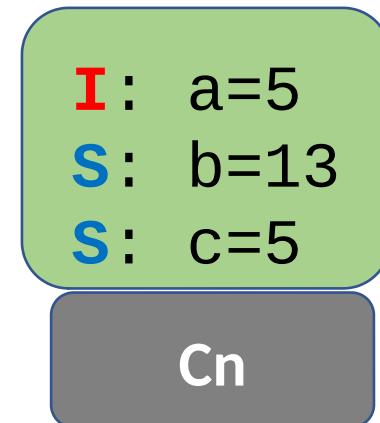
MSI protocol is a simple cache coherence protocol.
In this protocol, each cache line is labeled with a state:

- **M**: cache block has been modified.
- **S**: other caches may be sharing this block.
- **I**: cache block is invalid

If Cn wants to load 'a' then 'a' will be brought from the shared memory to the local cache of Cn.



...



Programming multicore platforms

- Option 1: Program directly targeting processor cores
 - Programmer takes care of synchronization
 - Painful and error-prone
- Option 2: Use a *concurrency platform*
 - It abstracts processor cores, handles synchronization and communication protocols, and performs load balancing.
 - Hence offers much easier multicore programming environment
 - Examples:
 - **Pthreads** and WinAPI threads
 - OpenMP