

File systems

File System

- Function: main permanent data storage
Speed bottleneck!
- Capacity not a problem nowadays: 2 TB disks even for PC.
But **backup becoming a problem.**
- Logical view (view of programmer): **tree structure of files** together with read/write operation and creation of directories
- Physical view: **sequence of blocks**, which can be read and written. OS has to map logical view to physical view, must impose tree structure and assign blocks for each file

Two main possibilities to realize filesystem:

- **Linked list**: Each block contains pointer to next
⇒ Problem: random access (`seek()`) costly: have to go through whole file until desired position.
- **Indexed allocation**: Store pointers in one location: so-called index block (similar to page table). To cope with vastly differing file sizes, may introduce **indirect index blocks**.

Index blocks are called **inodes** in Unix.

Inodes store additional information about the file (eg size, permissions)


Worked Example

Worked example – based on
<http://www.tavi.co.uk/phobos/fat.html>

Example: FAT

- F(ile) A(llocation) T(able) – dates back to 70s.
- Useful for explaining filesystem concepts, modern filesystems are more complicated
- Variants FAT12, FAT16, FAT32 define number of bits per FAT entry – we focus on FAT16
- Sector = disk unit (e.g. 512 byte), aka block
- Cluster = multiple sectors (factor 1, 2, 4, ..., 128)
(here: assume cluster = 1 sector)
- Uses linked list (“cluster chain”) to group clusters

Example: FAT16 Structure



Example: FAT16 Bootsector

0x13a1 = 5025 sectors

Number of FAT copies → 01


Block 0 (0x0000)

0x200 = 512 byte/sector

	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f
000	eb	3c	90	49	42	4d	2d	37	2e	30	20	00	02	01	01	00
010	01	40	00	a1	13	f8	14	00	0a	00	01	00	00	00	00	00
020	00	00	00	00	00	00	29	2a	65	bc	00	43	4f	38	38	33
030	2d	41	32	20	20	20	46	41	54	31	36	20	20	20	fa	31
040	c0	8e	d0	bc	00	7c	fb	8e	d4	e8	00	00	5e	83	05	19
050	bb	07	00	fc	ac	84	c0	74	06	b4	0e	cd	10	eb	f5	30
060	e4	cd	16	cd	19	0d	0a	4e	6f	6e	2d	73	79	73	74	65
070	6d	20	64	69	73	6b	0d	0a	50	72	65	73	73	20	61	6e
080	79	20	6b	65	79	20	74	6f	20	72	65	62	6f	6f	74	0d
090	0a	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00

Filesystem type Start of bootloader

Example: FAT16 with Offsets



Example: File Allocation Table

FAT_begin

FFF0	FFFF	0003	2	0004	3
0006	4	0000	5	FFFF	6
0009	8	FFFF	9	0000	A
000F	C	0010	D	0000	E
FFFF	10	0000	12	0000	13
				0000	14

2 → 3 → 4 → 6

7 → 8 → 9

C → F → D → 10

Example: File in Root Directory

Block 21 (0x0015)

	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f
000	43	4f	38	38	33	2d	41	32	20	20	20	28	00	00	00	C0883-A2 (....
010	00	00	00	00	00	00	91	9e	65	39	00	00	00	00	00e9.....
020	46	4f	4f	42	41	52	20	20	54	58	54	21	00	a3	91	9e FOOBAR TXT!....
030	65	39	65	39	00	00	91	9e	65	39	c6	10	1a	00	00	e9e9....e9.....
040	4e	45	54	57	4f	52	4b	20	56	52	53	20	00	b6	91	9e NETWORK VRS
050	65	39	65	39	00	00	91	9e	65	39	4e	0f	92	06	00	00 e9e9....e9N.....
060	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00

sector 0xF4E

length 0x692 = 1682 byte

filename & extension

Example: File in FAT

Block 16 (0x0010)

	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f
000	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
010	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
020	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
030	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
040	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
050	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
060	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
070	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
080	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
090	00	00	00	00	00	00	00	00	00	00	00	4f	0f	50	0f	4 block cluster chain (2048 byte)
0a0	51	0f	ff	ff	00	00	00	00	00	00	00	00	00	00	00	00
0b0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00

Example: FAT Limits

- Max. volume size: 2 GB ($2^{16} \cdot 32\text{ kB}$)
- Max. file size: 2 GB
- Max. number of files: 65,460 (32 kB clusters)
- Max. filename length: 8 + 3
- FAT32 / exFAT have higher limits
- Newer filesystems (NTFS, ext4) also overcome these limits, using other data structures (e.g. B-tree for dir structure, bitmap for allocation)

Further Aspects

Further aspects of filesystems

Caching

Disk blocks used for storing directories or recently used files cached in main memory

Blocks periodically written to disk

⇒ Big efficiency gain

Inconsistency arises when system crashes

Reason why computers must be shutdown properly

Journaling File Systems

To minimise data loss at system crashes, ideas from databases are used:

- Define **Transaction points**: Points where cache is written to disk
 - ⇒ Have consistent state
- Keep log-file for each write-operation
 - Log enough information to unravel any changes done after latest transaction point

Disk Access

Disk access contains three parts:

- **Seek**: head moves to appropriate track
- **Latency**: correct block is under head
- **Transfer**: data transfer

HDDs: Time necessary for seek and latency dwarfs transfer time
⇒ Distribution of data and scheduling algorithms have vital impact on performance for HDDs, less so for SSDs

Disk Scheduling Algorithms

Standard algorithms apply, adapted to the special situation:

- 1.) **FCFS**: easiest to implement, but: may require lots of head movements
- 2.) **Shortest Seek Time First**: Select job with minimal head movement

Problems:

- may cause starvation
- Tracks in the middle of disk preferred

Algorithm does not minimise number of head movements

- 3.) **SCAN-scheduling**: Head continuously scans the disk from en to end (lift strategy)
⇒ solves the fairness and starvation problem of SSTF

Improvement: **LOOK-scheduling**:
head only moved as far as last request (lift strategy).

Particular tasks may require different disk access algorithms

Example : Swap space management

Speed absolutely crucial ⇒ different treatment:

- Swap space stored on **separate partition**
- Indirect access methods not used
- **Special algorithms used for access of blocks**
Optimised for speed at the cost of space (eg increased internal fragmentation)

Linux Implementation

Interoperability with Windows and Mac requires support of different file systems (eg vfat)

⇒ Linux implements common interface for all filesystems

Common interface called **virtual file system**

virtual file system maintains

- inodes for files and directories
- caches, in particular for directories
- superblocks for file systems

All system calls (eg open, read, write and close) first go to virtual file system

If necessary, virtual file system selects appropriate operation from real file system

Disk Scheduler

Kernel makes it possible to have different schedulers for different file systems

Default scheduler (Completely Fair Queuing) based on lift strategy
have in addition separate queue for disk requests for each process
queues served in Round-Robin fashion

Have in addition No-op scheduler: implements FIFO

Suitable for SSD's where access time for all sectors is equal